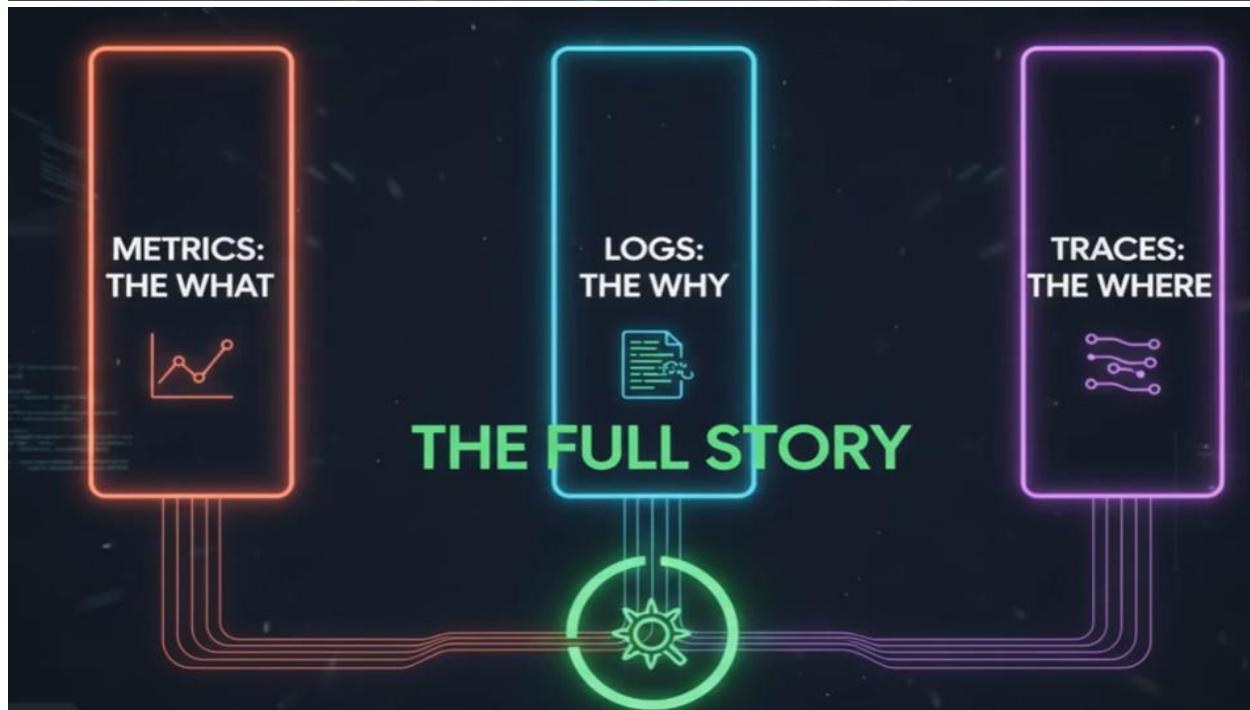
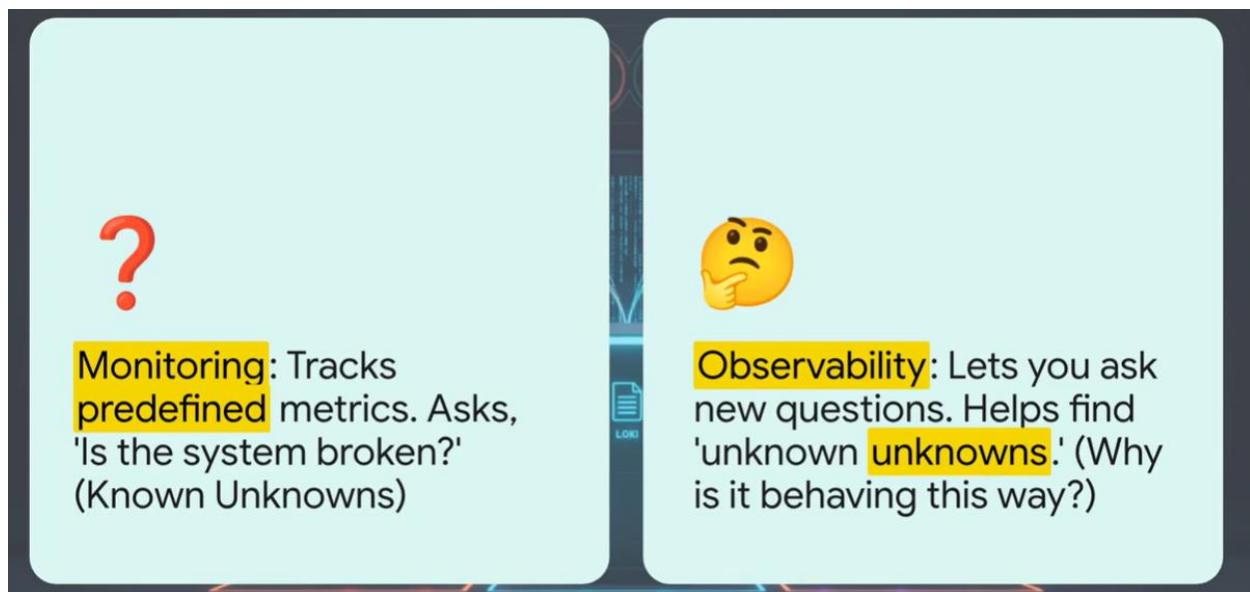


The Grafana LGTM Stack

- **Loki:** for logs
- **Grafana:** for visualization
- **Tempo:** for traces
- **Mimir / Prometheus:** for metrics



Metrics

Metrics tell you what is happening

Logs:

Logs tell you why is happening

Traces:

Traces tell you where is happening

```
Terminal - observability-lab
# What You'll Have in 5 Minutes:
[ 3-Node Kubernetes Cluster
  └─ 1 Control Plane (the brain)
      └─ 2 Worker Nodes (run your apps)
          └─ Helm (package manager)
              └─ Ingress (traffic routing)
]
# PLUS: Live nginx deployment demo!
```

THE LGTM STACK

PROMETHEUS

Metrics Engine

GRAFANA

Visualization

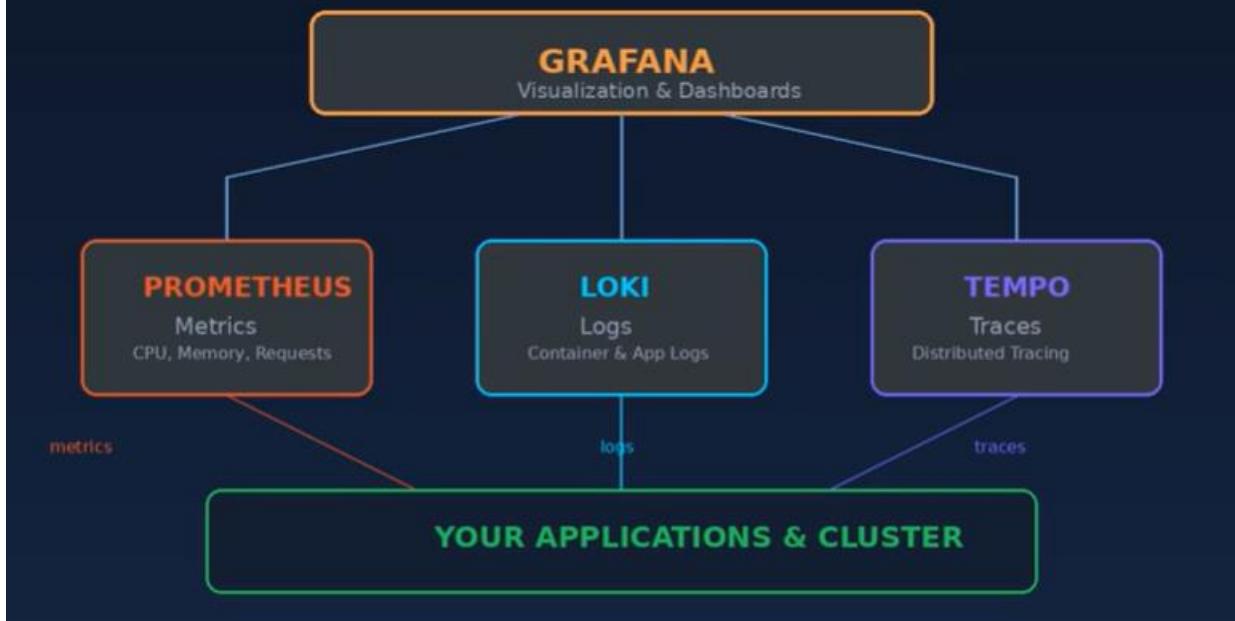
LOKI

Log Aggregation

TEMPO

Distributed Tracing

LGTM Stack Architecture



What You Built!

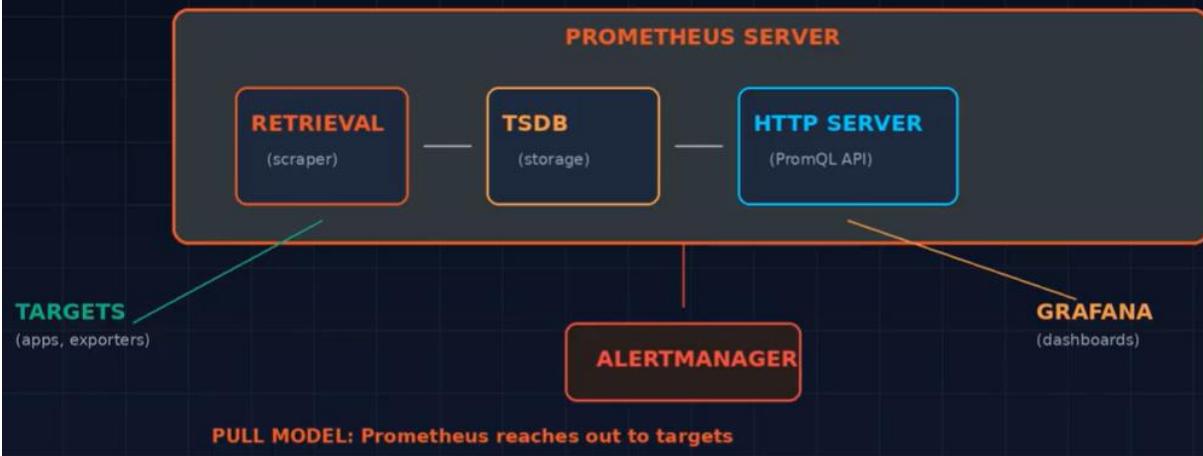
- ✓ **Prometheus**
Scraping metrics every 15 seconds
- ✓ **Grafana**
Pre-configured with all data sources
- ✓ **Loki**
Collecting logs from every pod
- ✓ **Tempo**
Ready for distributed traces



PROMETHEUS

How Metrics Collection REALLY Works

Prometheus Architecture



PULL vs PUSH



Use PUSHGATEWAY only for short-lived batch jobs

COUNTER

Only goes UP (or resets to zero)

Examples:

- http_requests_total
- errors_total
- bytes_sent_total

□ Use rate() - never raw value!



Only UP (resets on restart)

GAUGE

Goes UP and DOWN

Examples:

- memory_usage_bytes
- cpu_percent
- queue_depth

□ Current value is meaningful



Goes UP and DOWN

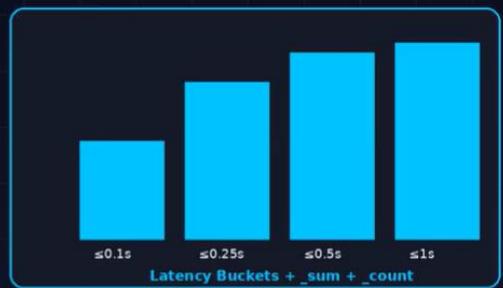
HISTOGRAM

Counts values in BUCKETS

Examples:

- request_duration_seconds_bucket
- _sum
- _count

□ Use histogram_quantile() for percentiles



Latency Buckets + _sum + _count

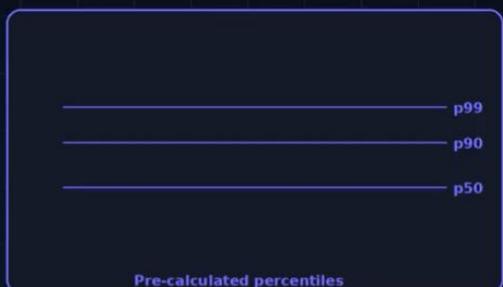
SUMMARY

Pre-calculated percentiles (client-side)

Examples:

- gc_duration_seconds{quantile="0.99"}

□ Prefer histogram for new metrics



Pre-calculated percentiles

Metric Naming Conventions

- **snake_case:** lowercase with underscores
- **Include units:** _seconds, _bytes, _total
- **Namespace prefix:** prometheus_, node_, http_

□ Good Examples:

```
http_requests_total  
node_memory_Active_bytes  
request_duration_seconds
```

□ Bad Examples:

```
httpRequests (wrong case)  
request_time (no unit)  
memory (too vague)
```

Labels = Dimensional Data

One metric name, infinite dimensions:

```
http_requests_total{method="GET", status="200"} 15234  
http_requests_total{method="GET", status="500"} 23  
http_requests_total{method="POST", status="201"} 890
```

Each unique label combination = separate time series

⚠ HIGH CARDINALITY WARNING

NEVER use as labels: user_id, request_id, trace_id, or unbounded values

1 million users = 1 million time series = dead Prometheus!

The screenshot shows the Prometheus web interface at localhost:9090. The top navigation bar includes links for Alerts, Graph, and Status. The main content area is titled "Status → Targets". It displays a table with four rows of target information:

Endpoint	State	Labels	Last Scrape
prometheus-server:9090/metrics	UP	job="prometheus"	2s ago
node-exporter:9100/metrics	UP	job="node"	5s ago
kube-state-metrics:8080/metrics	UP	job="kube-state"	3s ago

Below the table, a message states: "All targets healthy = Prometheus is collecting data".

localhost:9090

PROMETHEUS

Graph → Query

```
prometheus_http_requests_total
```

Execute

Results (showing 4 of 12 series):

```
{handlers="/api/v1/query", code="200"} 8921
{handlers="/api/v1/targets", code="200"} 342
{handlers="/metrics", code="200"} 15234
```

See the labels? Each combination is tracked separately

localhost:9090

PROMETHEUS

Status → TSDB Status

Head Series:	45,231	Active time series
Head Chunks:	182,924	Data chunks in memory
WAL Size:	128 MB	Write-ahead log
Retention:	15 days	Data kept before deletion

Watch Head Series - If it explodes, you have cardinality issues

Kubernetes Service Discovery

Pod Annotations:

```
metadata:
  annotations:
    prometheus.io/scrape: "true"
    prometheus.io/port: "8080"
    prometheus.io/path: "/metrics"
```

Auto-Discovery Flow:

```
graph TD; A[Pod Created] --> B[K8s API Notifies]; B --> C[Prometheus Scrapes]
```

No config file changes needed!
Pods come and go, Prometheus adapts automatically

Critical Settings

SCRAPE INTERVAL

Default: 15 seconds

Lower = more data, more storage

Recommended: 15-30 seconds

RETENTION

Default: 15 days

Older data deleted automatically

Long-term: Use Thanos or Mimir

Helm values.yaml:

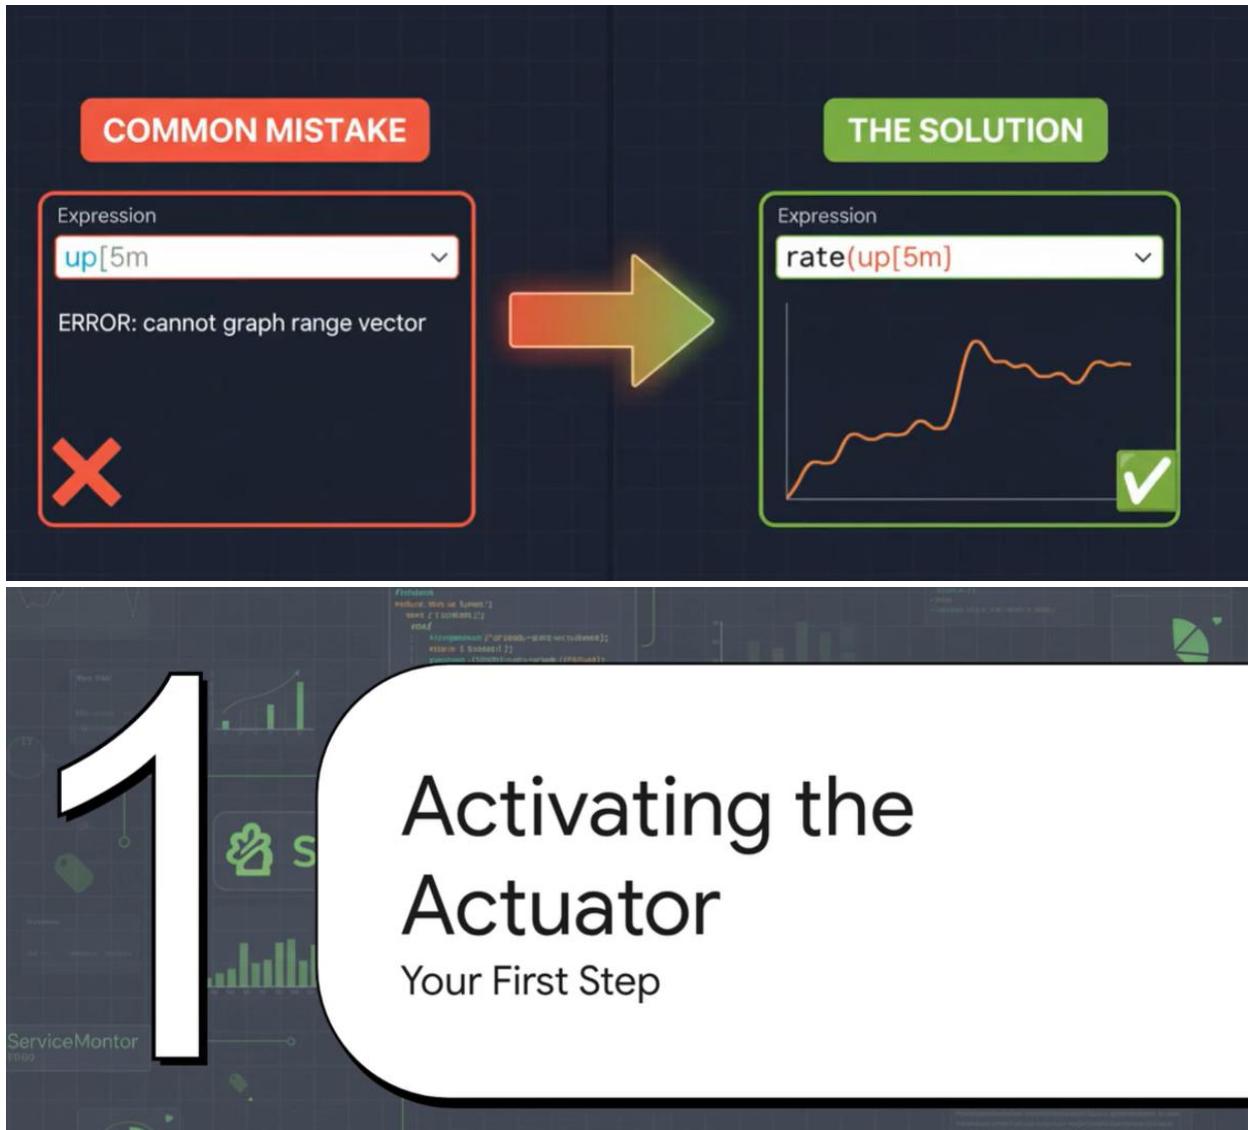
```
server.global.scrape_interval: 15s  
server.retention: 30d
```

“ Raw data is **useless** until you can **QUERY** it. Master the language behind every dashboard and alert.



Four Ways to Filter

- `=`: Exact match (e.g., `job="prometheus"`)
- `!=`: Not equal (e.g., `job!="test"`)
- `=~`: Regex match (e.g., `job=~"prod.*"`)
- `!~`: Negative regex match (e.g., `job!~".*test.*"`)



1

Activating the Actuator

Your First Step

Getting Started

PROMETHEUS & METRICS

6/12

1. Add Dependencies

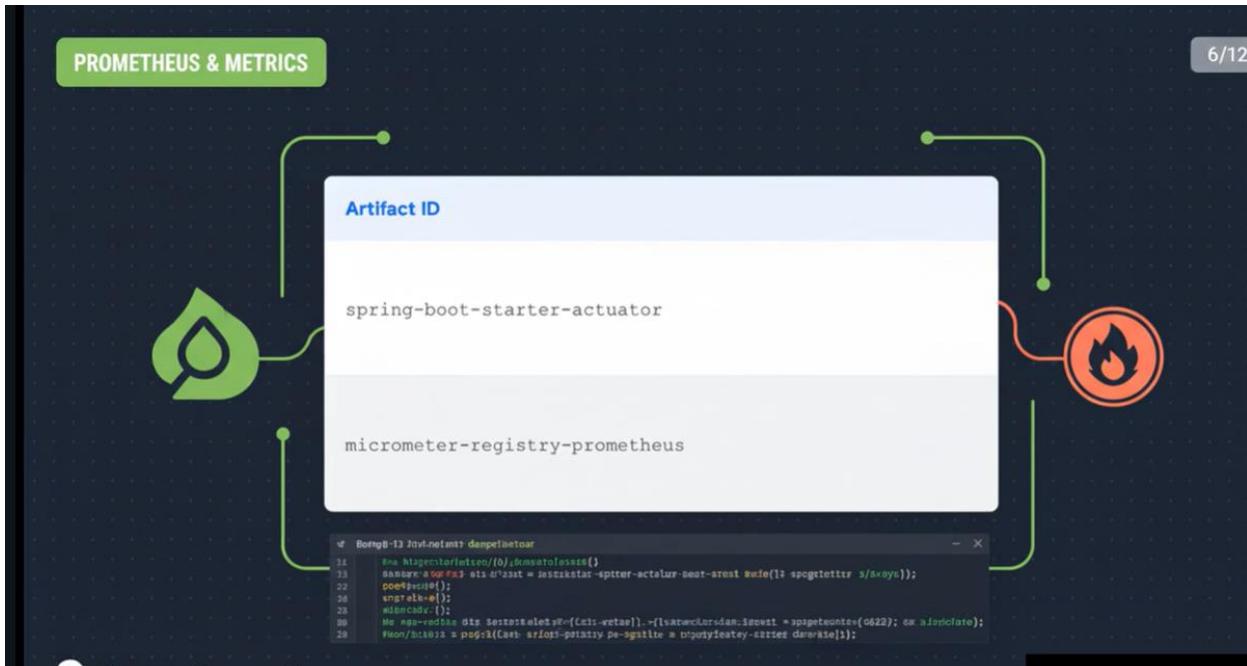
Include Actuator and the Prometheus Micrometer registry.

2. Enable Endpoint

Expose the Prometheus endpoint via configuration registry.

3. Access Metrics

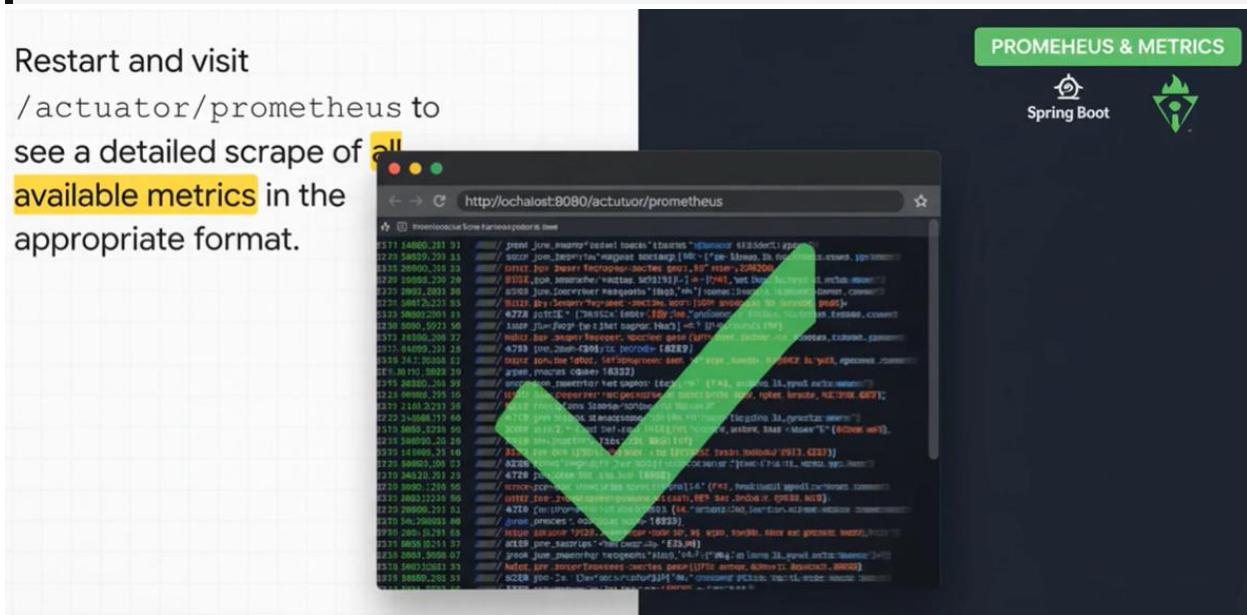
Restart and view the new metrics endpoint.



`dpoints.web.exposure.include=heal
tion to expose web endpoints.`

Restart and visit

/actuator/prometheus to see a detailed scrape of all available metrics in the appropriate format.



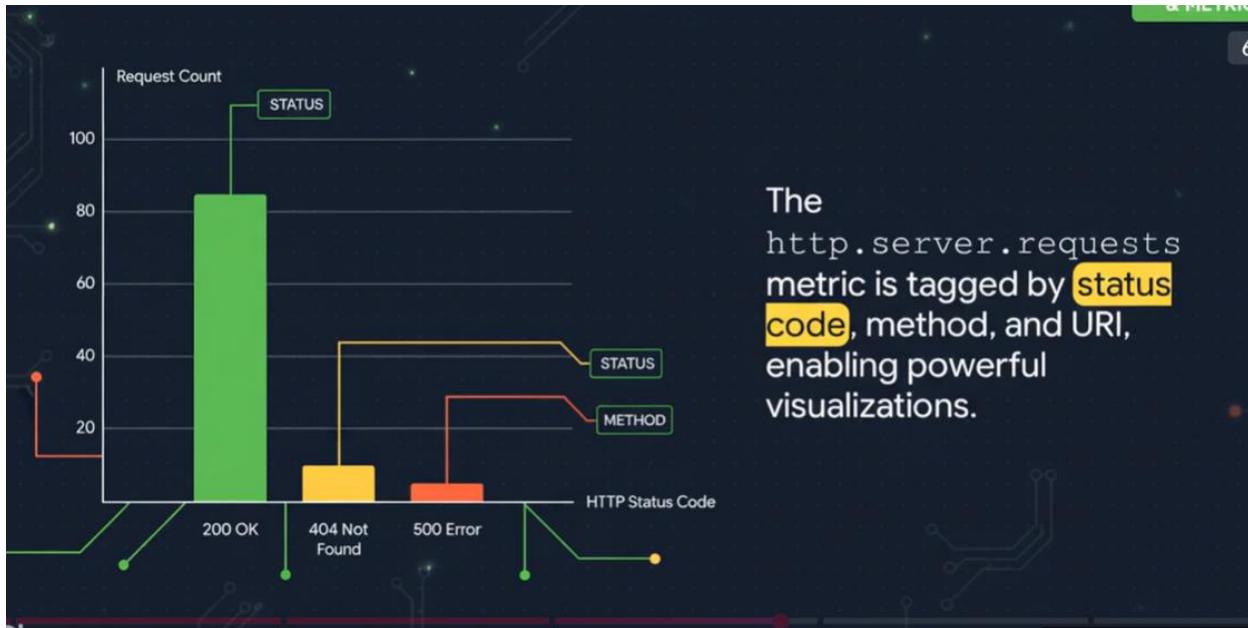


PROMETHEUS & METRICS | 6/12

Java

Metric Prefix	Description
jvm.	Memory, GC, and thread utilization
process.	CPU usage and uptime metrics
http.server.requests	Latency for Spring MVC requests

Gheware

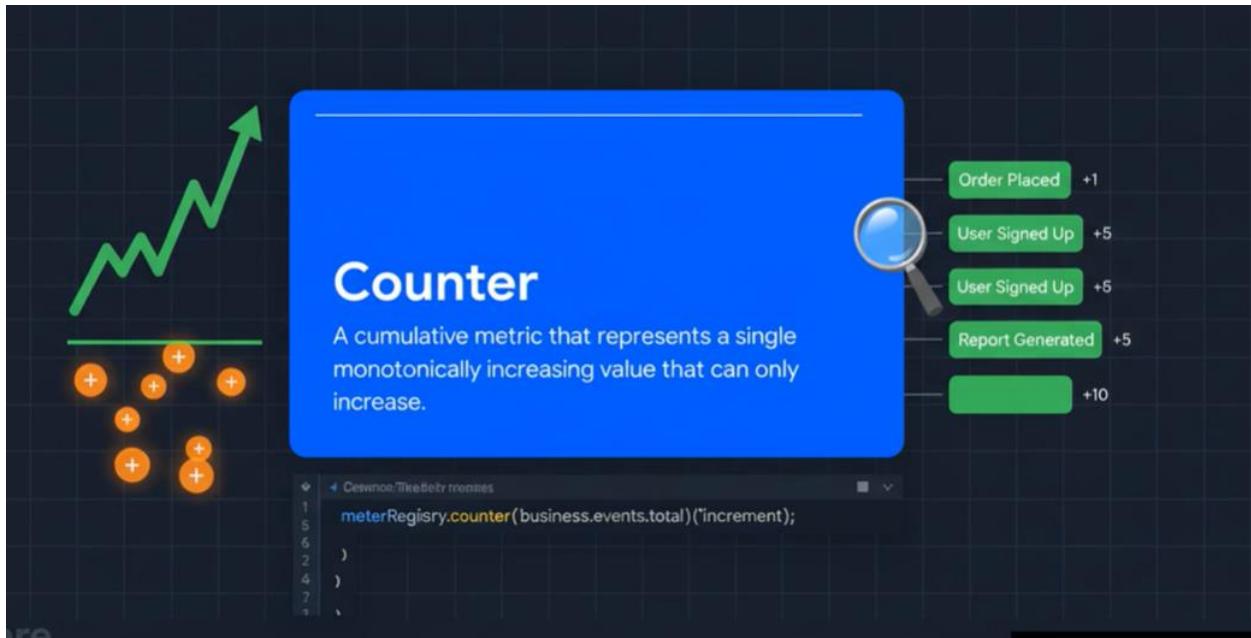


PROMETHEUS & METRICS 6/12

Custom Business Metrics

Beyond the Technical

Gheware DevOps AI



To register custom metrics,
inject `MeterRegistry` into
your component, then build
and increment your counter.

Gauge

A metric that represents a single numerical value that can arbitrarily go up and down over time.

PROMETHEUS & METRICS

PROMETHEUS & METRICS 6/12

Prometheus & Grafana
Full Visibility

Kubernetes Discovery

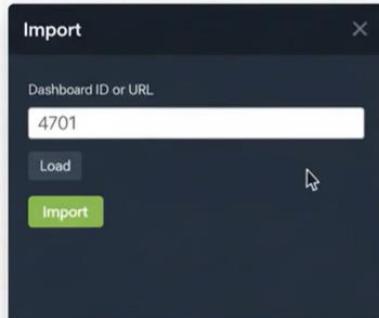


```
Wiers-Pond
apiVersion: V
1  assertion: Pod
2    himd
3      my.spring-Boot-app
4    annotations:
5      annotations:
6        prometheus.io/path: "/actuator/prometheus"
7        prometheus.io/path: "/actuator: "8880"
8      containers:
9        containers: ...
10     );
11   );
12 );
```

Auto-Discovery



Use a prebuilt dashboard for Micrometer instrumented applications to **get started instantly** without custom configuration.



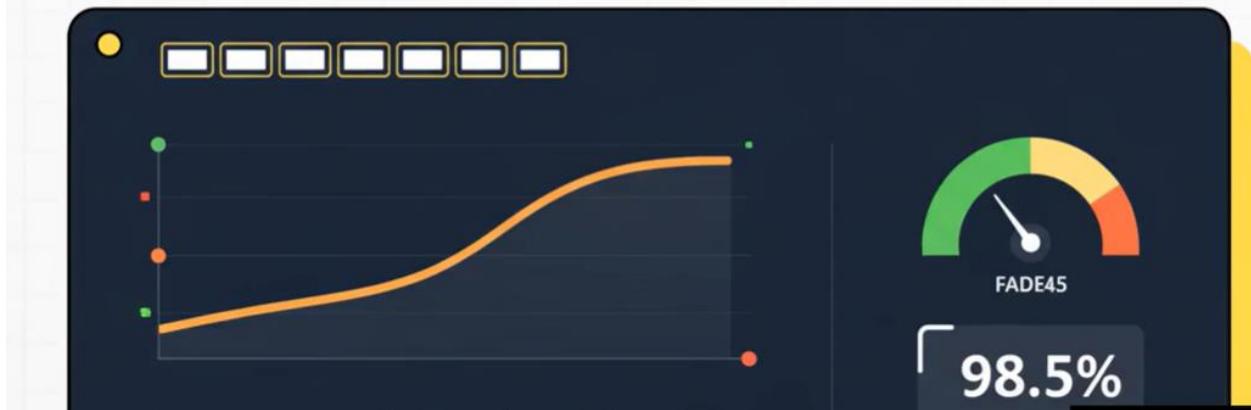
PROMHEUS & METRICS



Instantly visualize key performance indicators like JVM Thread States, Garbage Collection, CPU Usage, and more.



Production Grafana Dashboards



01

Dashboard
Design

02

Choosing The
Right Panel

03

Dynamic
Dashboards

04

Organizing For
Clarity

05

Context w/
Annotations

06

Don't Reinvent
The Wheel

1

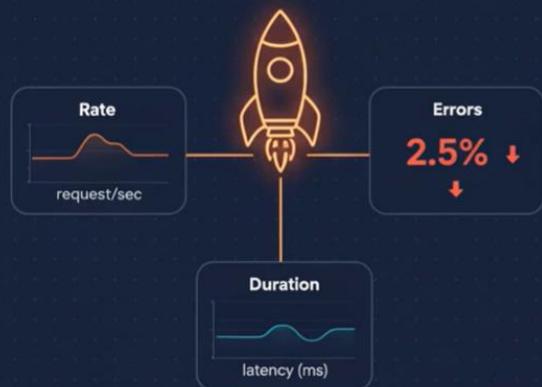
Dashboard Design

Start with a framework



RED Method

For Applications & Services



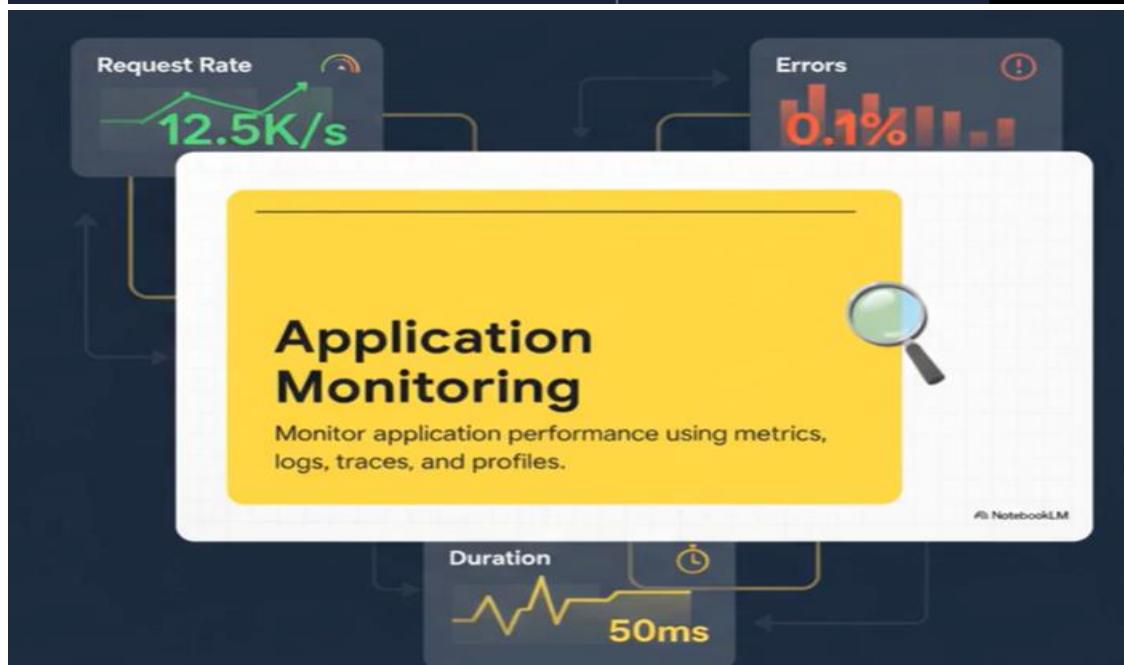
USE Method

For Infrastructure & Resources



Application Monitoring

Monitor application performance using metrics, logs, traces, and profiles.

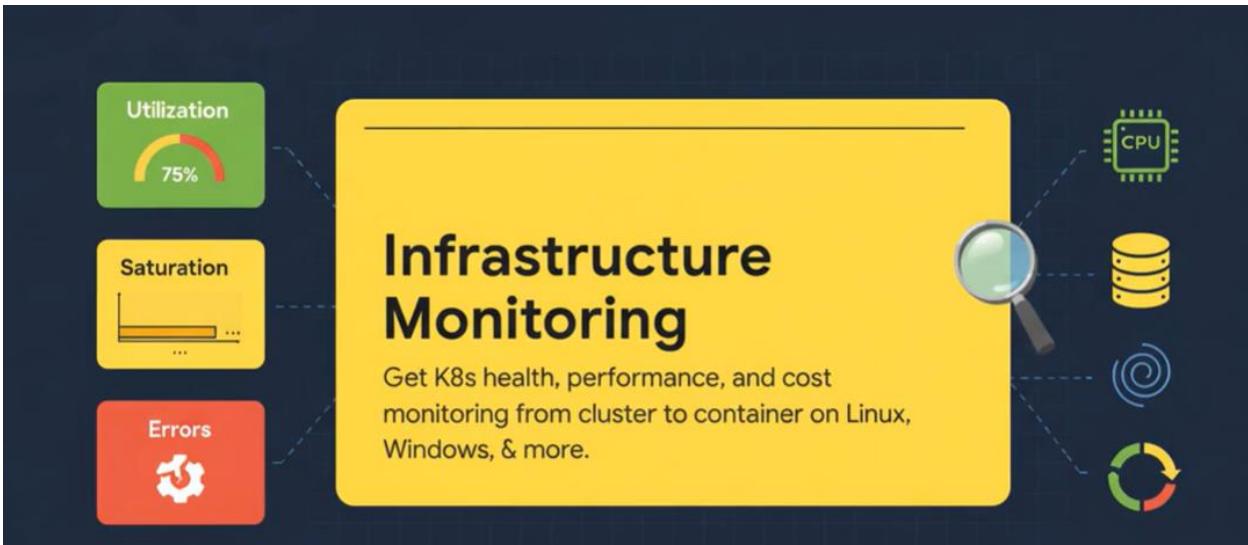


Any service we are running need to know three things

Rate: How much traffic are we getting

Errors: How much traffic is failing

Duration: How long is it taken

A large white number "2" is overlaid on a background of various data visualization panels, including time series graphs and pie charts. A white callout bubble with a black border contains the text "Choosing The Right Panel" and "Visualize your data".

2

Choosing The Right Panel

Visualize your data

Essential Panel Types



Time Series



Stat



Gauge

Name	Value	Status	Status
1	65	155	13
2	43	185	53
FluxDensity	565	28%	Red

Table



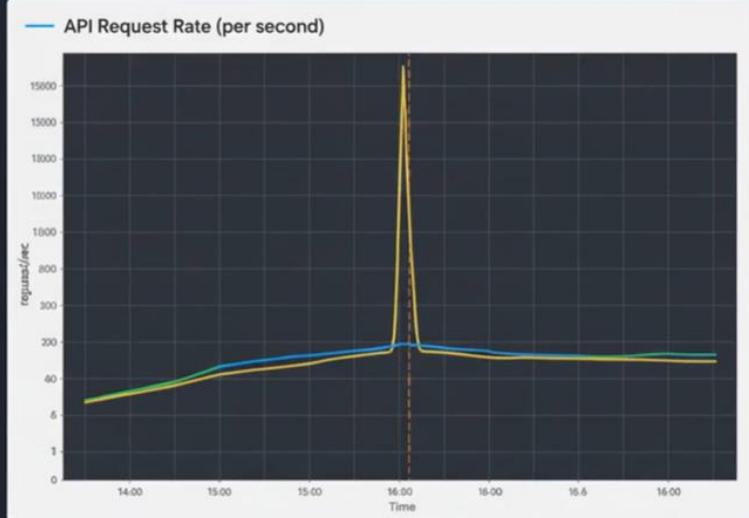
Heatmap



Bar Gauge

Time Series

The default choice for tracking metrics over time.



Stat Panel: At-a-glance KPIs that demand attention.

Availability SLO

99.98%

Requests (24hr)

2.45M

Error Rate

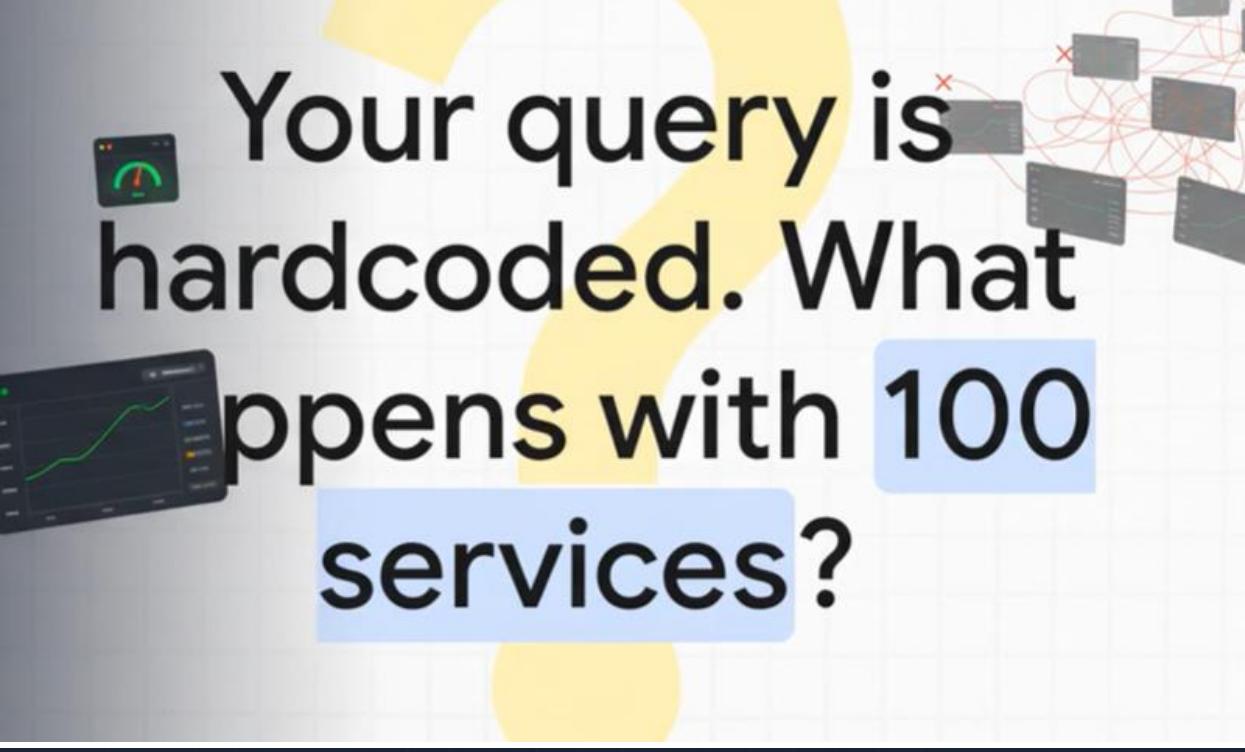
1.7%

Disk Space Utilization (/data)



The **Table** panel is perfect for displaying detailed, multi-dimensional data clearly.





Your query is hardcoded. What happens with 100 services?



Hardcoded
Query

```
rate(http_requesttotal{  
pod="payments-api-7f8C9d[5m)"}  
rate(http_requesttotal{  
pod="payments-api-7f8C9d[5m)"}  
rate(http_requesttotal{  
$pod"
```



Dynamic Query
with Variables

Pod
payments-api-5h4339
users-service-5h4j3k
inventory-db-8p9q7r

```
rate(http_requesttotal{  
$pod"}
```

Chaining Variables

GRAFANA DASHBOARDS 7

Get Namespaces

Create a variable that queries your data source for a list of items.

Get Pods

Create a nested variable that uses the first variable in its query.

Instant Update

Changing the variable value updates the data in your dashboard.

4

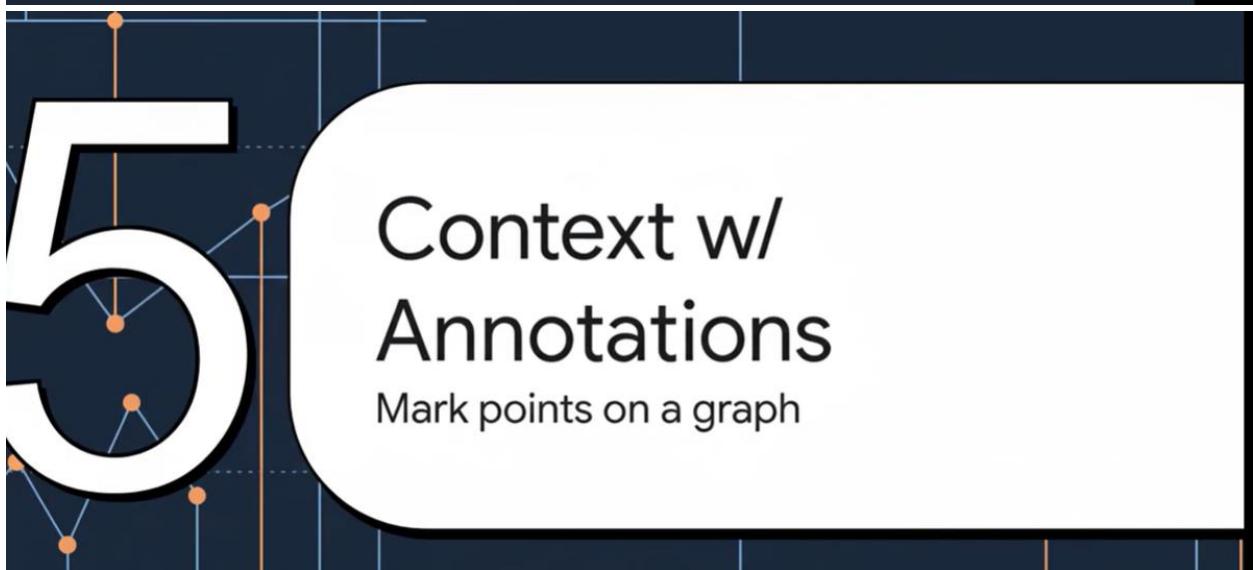
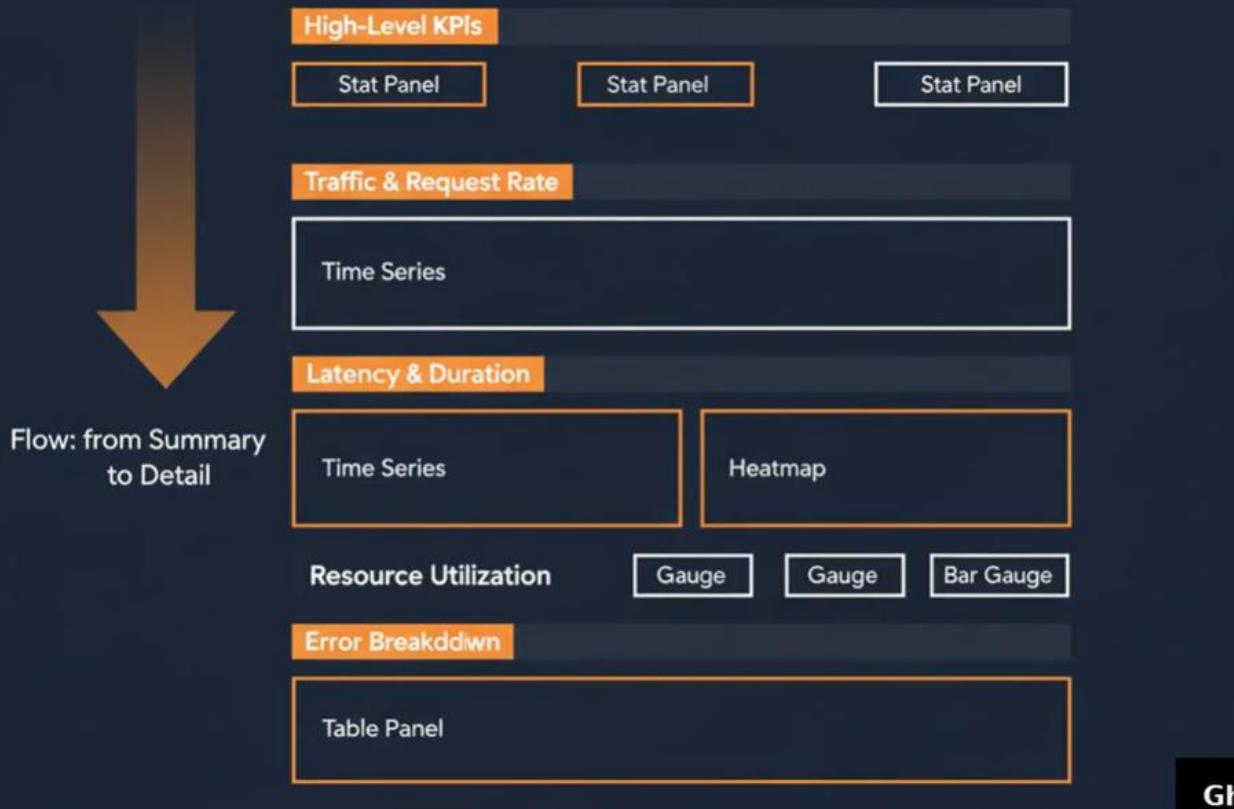
Organizing For Clarity

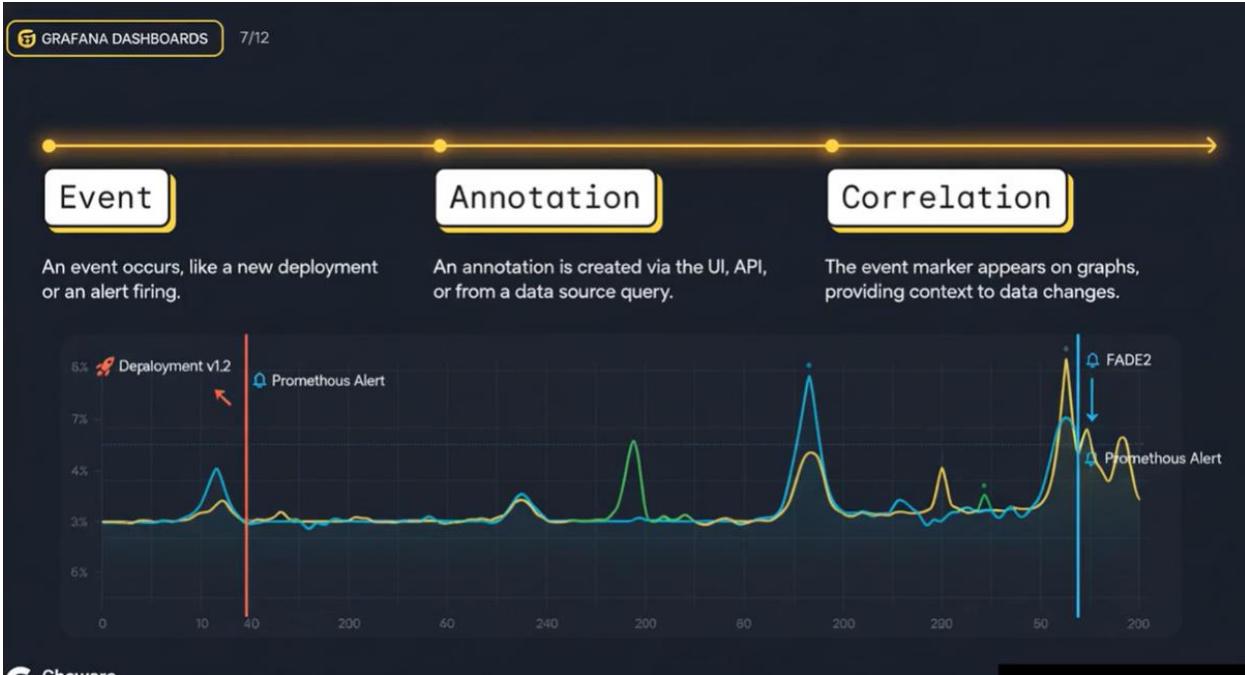
Structure into rows

7/12

Grafana

Dashboard Layout Blueprint





Name	Use Case
Kubernetes Monitoring	K8s health & performance
MongoDB Data	Visualize MongoDB data
Jira Data	Visualize Jira data
InfluxDB Dashboards	Monitor telegraf metrics

DOs

 Start with a Methodology

 Use Dynamic Variables

 Organize in Logical Rows

 Treat Dashboards as Code



DON'T

 Cram Everything In

 ~~server01~~

 Hardcode Values

 Use Random Graphs



You can visualize data...
but how do you get
alerted?

Alerting: Prometheus &
Grafana

ALERTING 8/12

SYSTEM HEALTH: OK

Your dashboards are green,
but are you watching 24/7?

Let's build alerts that wake you
only when it truly matters

ALERTING 8/12

1

The Alerting Paradox

Fatigue vs. Action

Gheware DevOps AI

TOO MUCH NOISE **FIRING**

JUST **Gheware DevOps AI**

This image is a composite of two parts. The top part shows a screenshot of a system health dashboard with multiple green 'OK' status indicators. Overlaid text reads: 'Your dashboards are green, but are you watching 24/7?' followed by 'Let's build alerts that wake you only when it truly matters'. To the right is an illustration of a person sitting in a beanbag chair, holding a coffee cup, with a computer monitor in the background. The bottom part is a slide titled 'The Alerting Paradox' with the subtitle 'Fatigue vs. Action'. It features a large number '1'. The slide is decorated with various icons: a red circle with a bell labeled 'ALERTING' and '8/12'; a large white '1'; a bell icon; a stack of cards labeled 'CRITICAL'; a stack of cards labeled 'CRITICAL 1/16'; several small fire/flame icons; a hand holding a pile of small bells; a hand holding a small fire; and a smartphone with a Slack notification icon. Text at the bottom right says 'TOO MUCH NOISE' and 'FIRING', with arrows pointing to the phone and the notification. A red button labeled 'JUST' is at the bottom left. The Gheware DevOps AI logo is at the bottom right.

Too Many Alerts



Too Many Alerts: Leads to **alert fatigue**, ignored notifications, and missed **critical** incidents.

Too Few Alerts



ALERTING

8/12

“Every page should require human intelligence to resolve.

The central part of the image features a white rectangular box containing a quote: “Every page should require human intelligence to resolve.” The quote is split into four yellow-highlighted segments: “Every page”, “should”, “require human”, and “intelligence to resolve”. To the left of the quote is a stylized icon of a blue cylinder with a brain inside, connected to a circuit board. To the right is a monitoring interface with binary code, a red warning triangle with a skull, a yellow warning triangle with an exclamation mark, two green checkmarks, and a mobile phone icon. The Grafana logo is visible. The bottom left corner of the central box contains the word “Gheware”.

The Three Golden Rules of Alerting

1. Alert on SYMPTOMS



High Latency



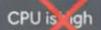
CPU Usage High

2. Every alert must be ACTIONABLE



CRITICAL: Service X Down
Down

[View Runbook](#)



CPU is high

3. Alerts must be URGENT



FIRING

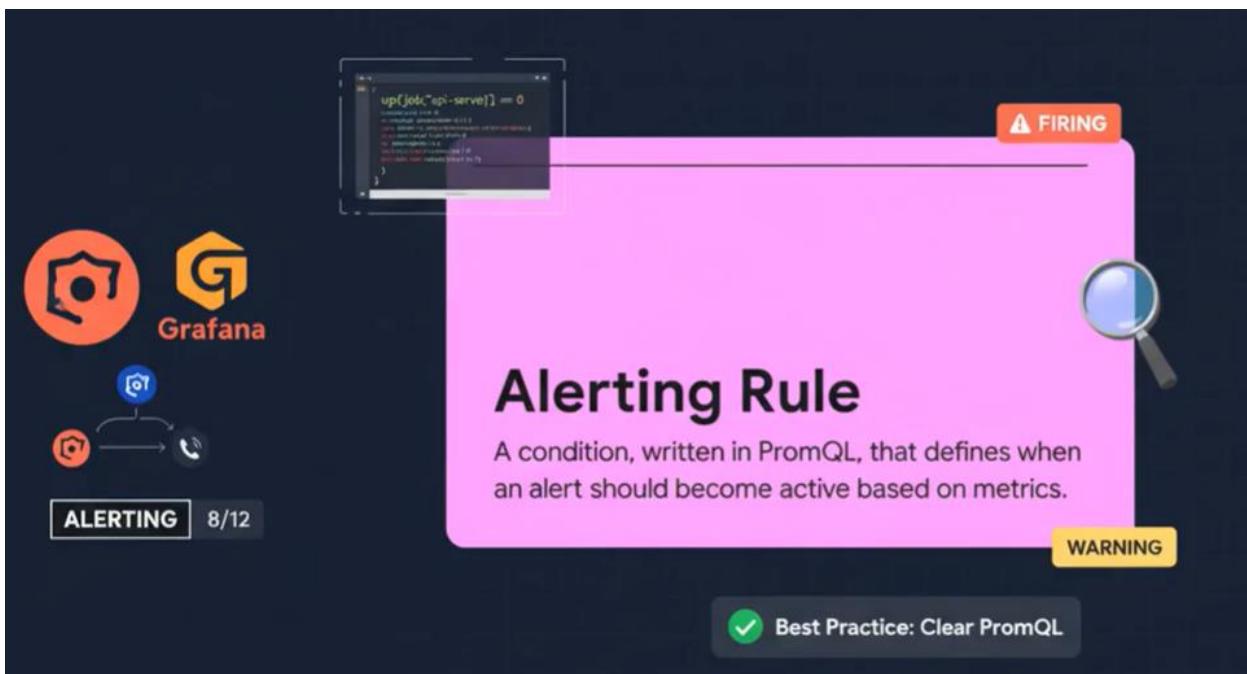
Service X Down



CPU is high

Prometheus's Engine

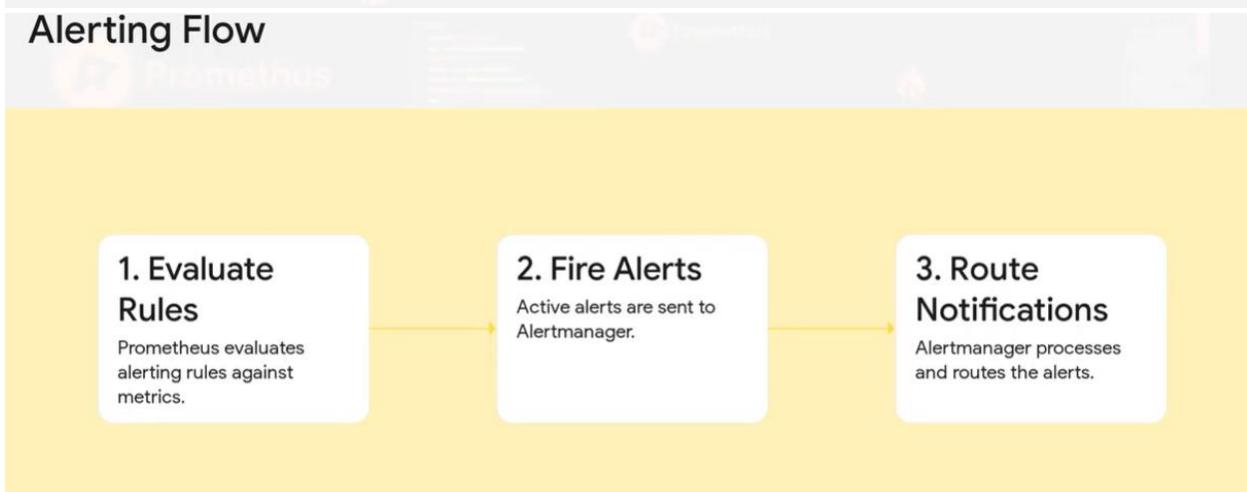
The brains of the op



Anatomy of an Alert Rule

- `expr`: The PromQL query defining the failure condition.
- `for`: Duration to prevent alerts from brief, flapping spikes.
- `labels`: Metadata for routing and severity (e.g., `severity: page`).
- `annotations`: Human-readable info, like a summary or runbook link.

Alerting Flow



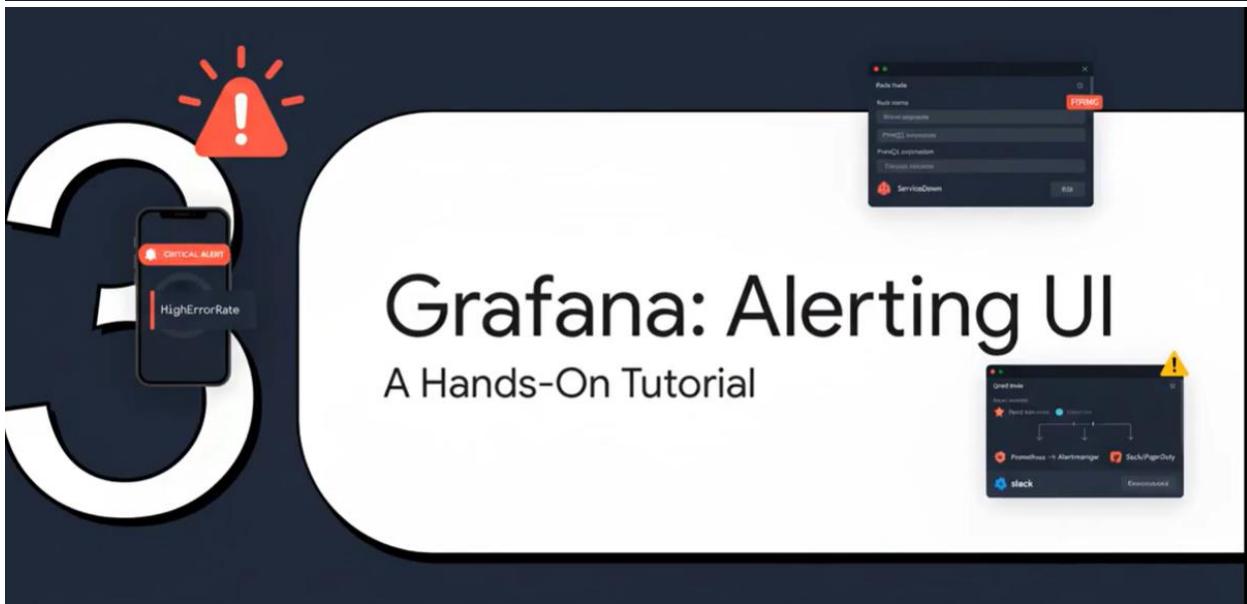
Alertmanager's Powers

The Traffic Cop.



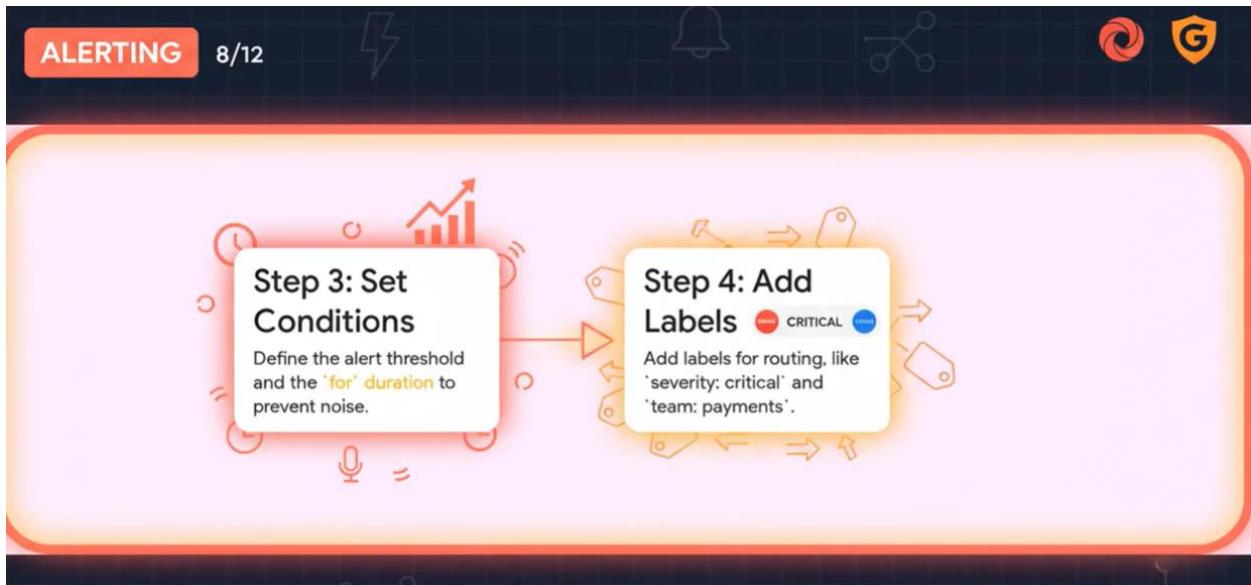
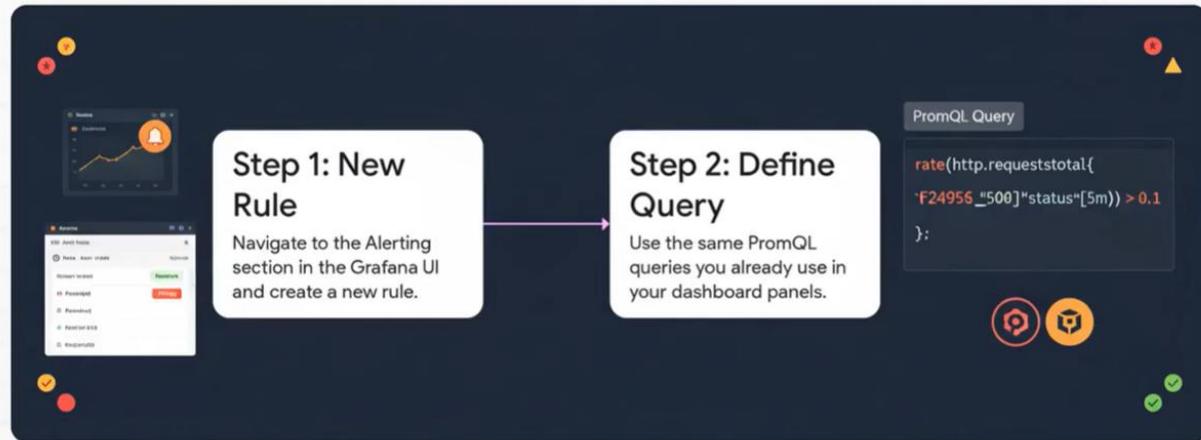
Grafana: Alerting UI

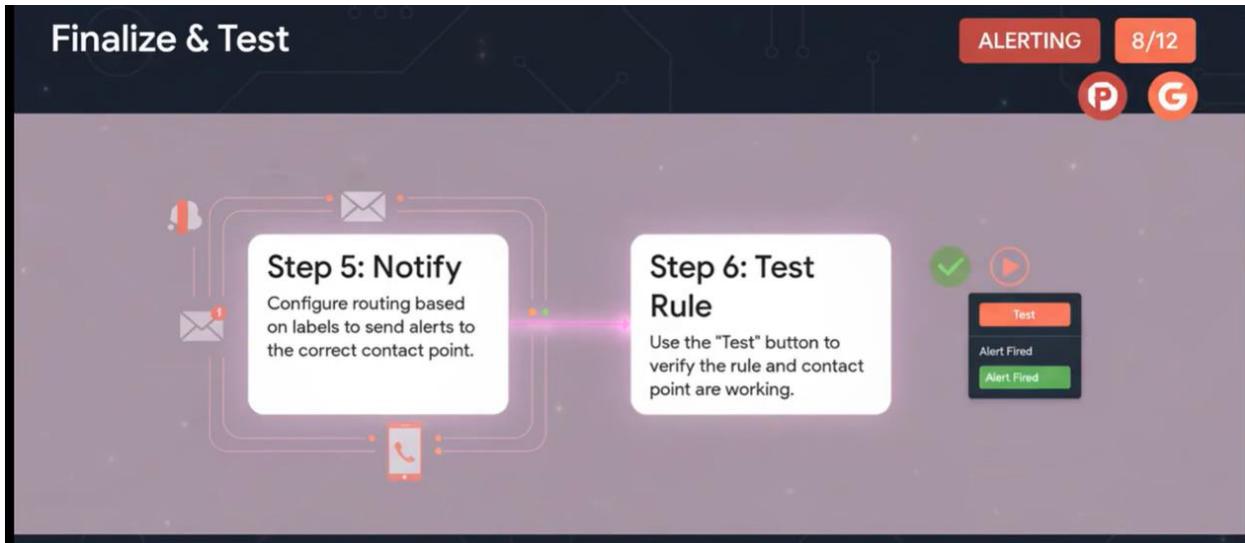
A Hands-On Tutorial



Create Alert in Grafana

ALERTING 8/12





The end result: a **clear**,
actionable notification in Slack
with a summary, description,
and **severity** to kickstart the
investigation. 

The image is a composite of several visual elements related to SRE and error budget management:

- Large Number 4:** A prominent white number 4 with a red outline.
- Alerting Flow Diagram:** Shows a process flow from "Alertmanger" (with a bell icon) to another alert manager instance, which then triggers "Receivers" (represented by a phone and a laptop).
- Grafana Unified Alerting Dashboard:** A screenshot of a Grafana dashboard titled "ALERTING". It shows an "Error Budget" chart with a green line decreasing over time. Key points on the chart include "SLO Target" (red dot), "Latency Exceeded" (yellow dot), and "Warning" (orange dot). The chart includes labels for "Remaining Budget" and "Burn Rate". The dashboard also displays "8/12" alerts, "5:54 / 7:17" time, and "Grafana Unified Alerting >".
- Error Budget Visualization:** A pink-bordered card titled "ALERTING" featuring the text "Error Budget" and "The acceptable amount of user-facing error". It includes a "SLO Target: 99.9%" indicator, a progress bar showing "Remaining Budget" (green) and "Burn Rate" (orange), and the text "43 minutes/month". A magnifying glass icon is positioned on the right side of the card.

 **DOS**

-  Include runbook links. 
-  **Include** runbook links. 
Use 'for' 'for' to prevent flapping.  **Group** alerts to reduce noise. 

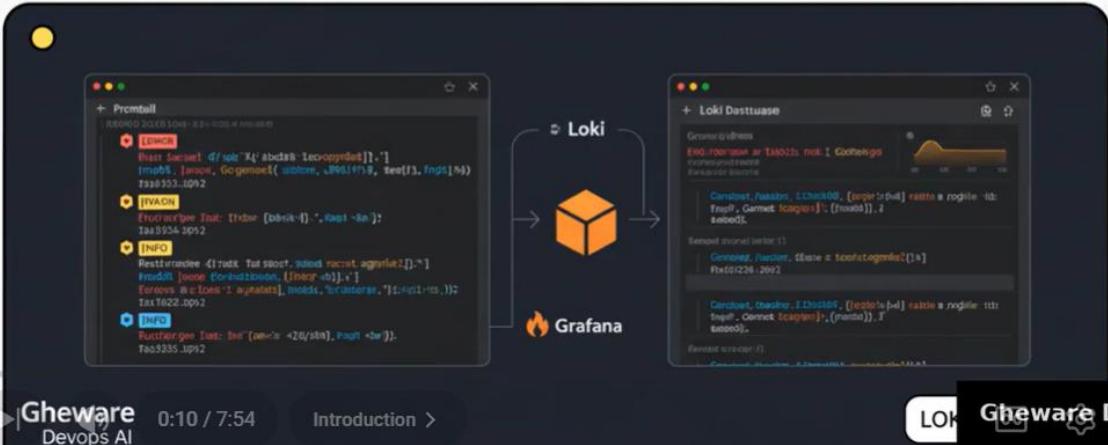
ALERTING 8/12

 **DON'TS**

-  Alert on every metric. 
-  **Alert** on every metric.
-  Forget 'for' duration. 
- Create alerts without runbooks. 



Chaos to Clarity with Loki



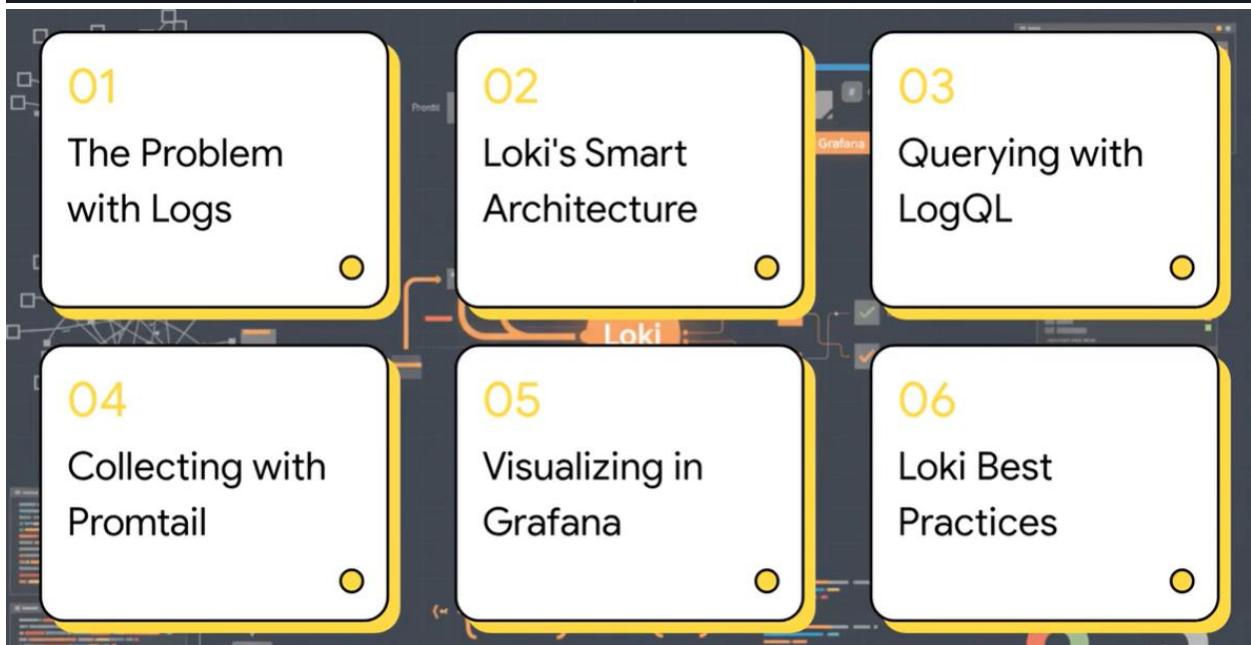
Alerts fired. You know
WHAT broke. But do
you know WHY?

Before: The Hunt



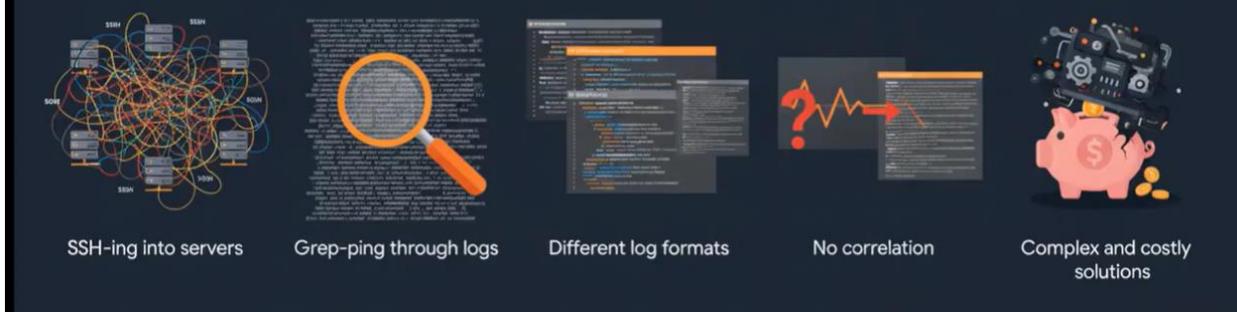
From many tools to
one unified view

After: The Promise of Loki

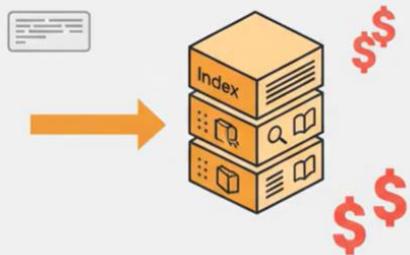




Distributed Logging Pain

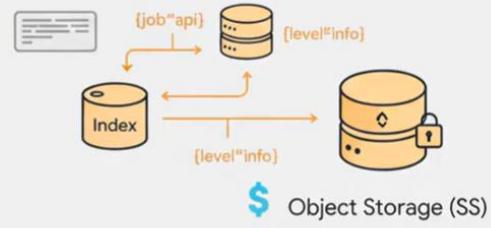


ELK Stack: Full-Text Indexing

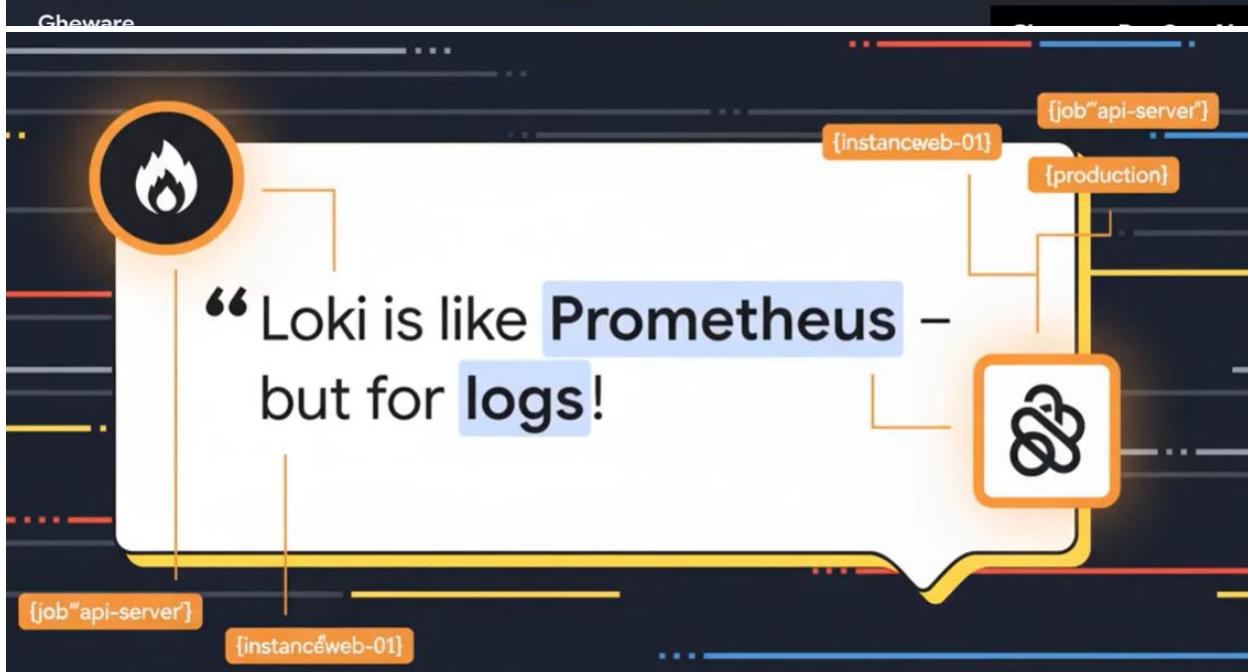


Indexes everything. High storage, powerful search.

Grafana Loki: Label Indexing



Indexes only labels. Low storage, cost-effective.



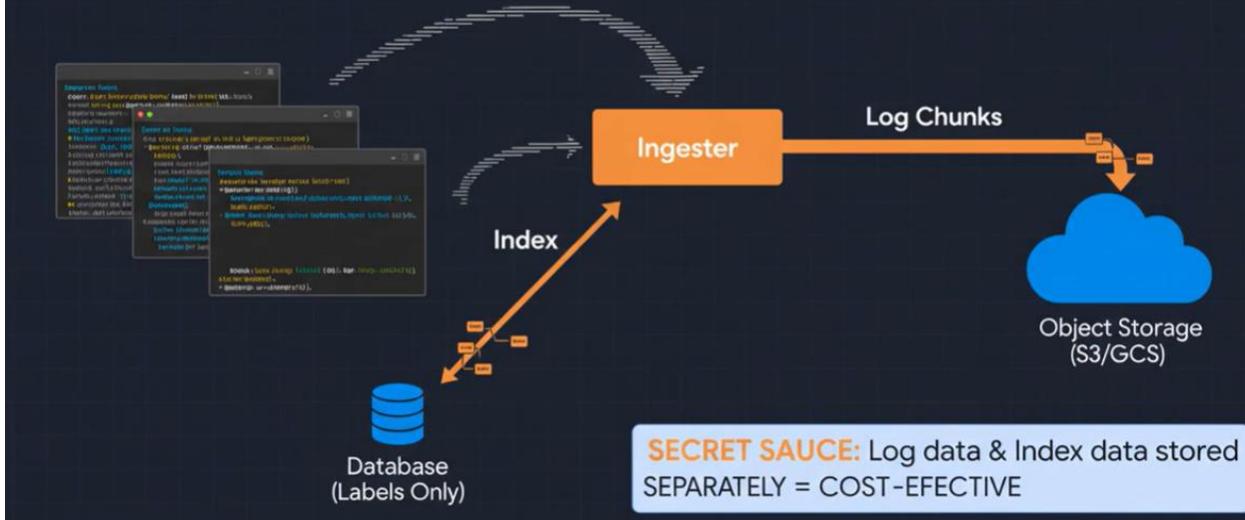


Loki's Data Flow

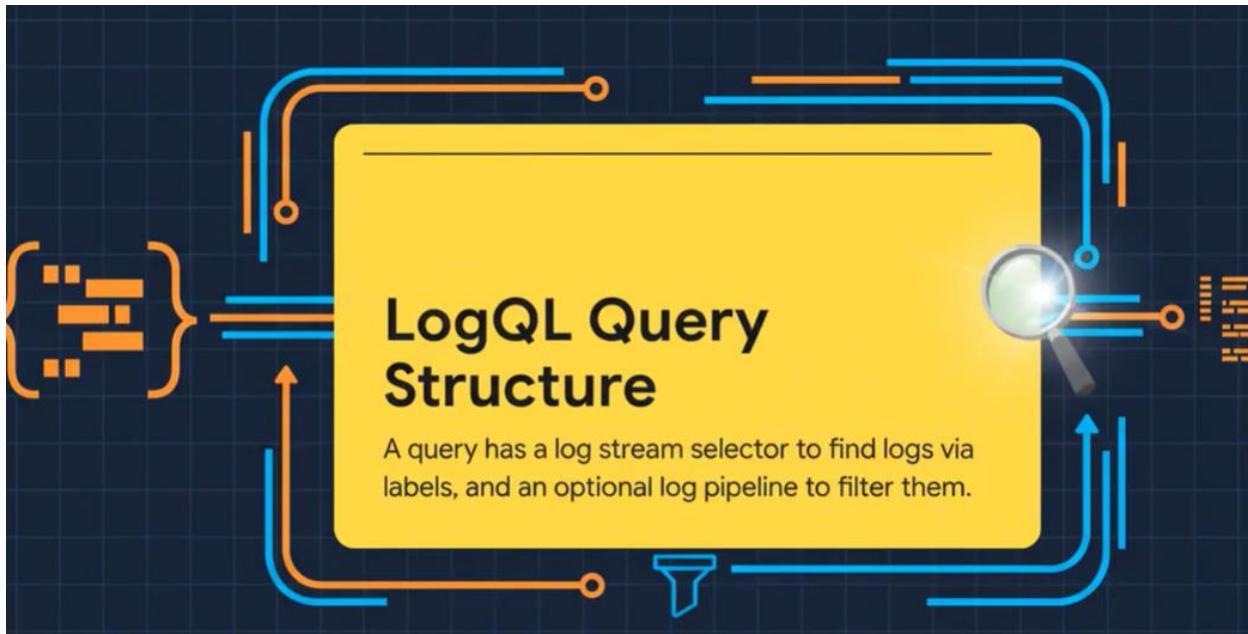


Loki Storage Explained

9/



The diagram shows a large number '3' on the left, with a terminal window displaying log entries. A line connects the terminal to a box containing a LogQL query: `{job="api"} l: cronjob {job="api"} l: "error"`. This query is connected via arrows to various monitoring tools: Promtail, Loki, and Grafana. The text "Querying with LogQL" and "Connecting to your logs" is overlaid on the right side of the diagram.



Operator	Description
=	Exactly equal
!=	Not equal
=~	Regex matches
!~	Regex does not match
*	=
*	~*

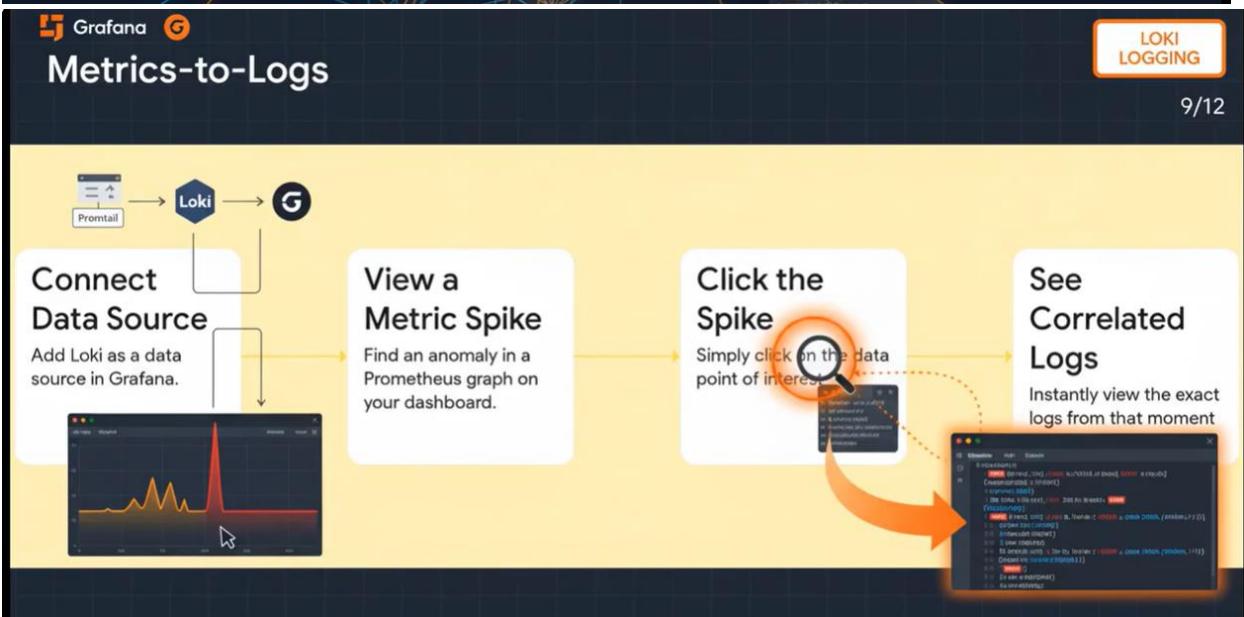
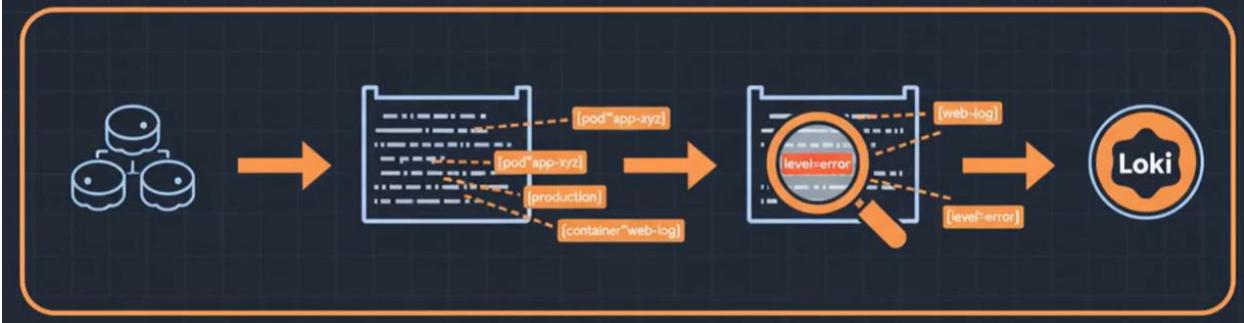
Beyond Basic Queries

- Parse logs on the fly with `| json` or `| logfmt`.
- Filter on newly extracted fields, like `duration > 500ms`.
- Create metrics from logs, just like in Prometheus!

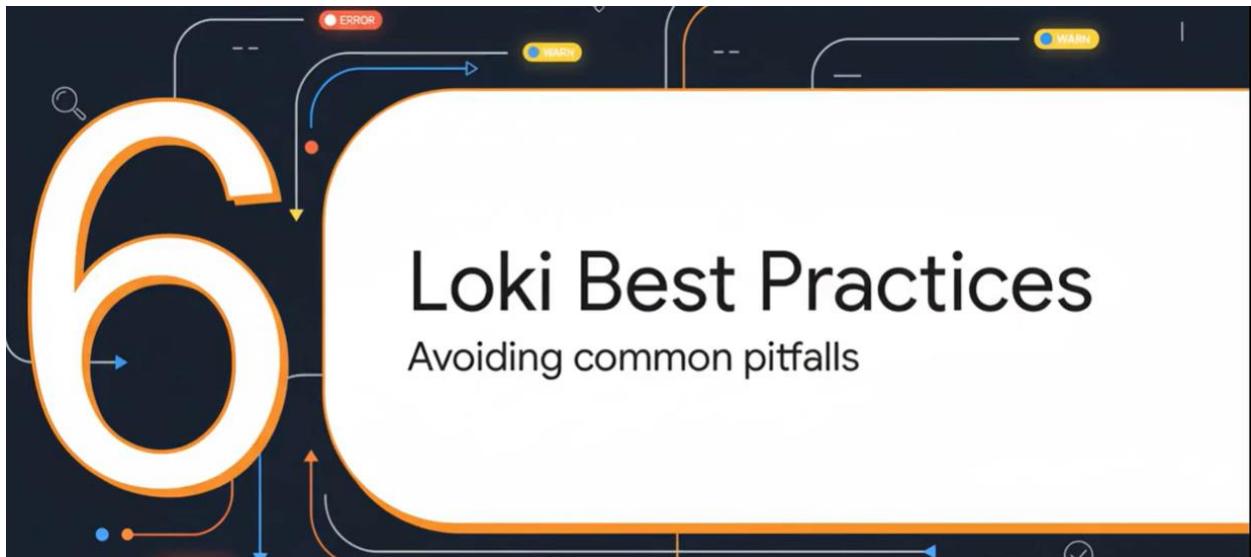


A simple YAML file tells Promtail where to send logs (clients) and what to scrape (`scrape_configs`), just like Prometheus.

Promtail's Pipeline



Your Prometheus metrics and Loki logs, side-by-side. Go from 'What broke?' to 'Why?' in seconds, without ever leaving your dashboard.



DO: Low Cardinality



index

{ job=api }

{ job=db }

{ env=prod }

DON'T: High Cardinality



Few, distinct streams → Efficient Index.

Chaotic, infinite streams → Kills performance.

Production Takeaways



Manage Cardinality

Keep label values limited and stable.



Use Structured Logs

JSON or logfmt makes parsing easy.



Optimize Queries

Start with restrictive label selectors.

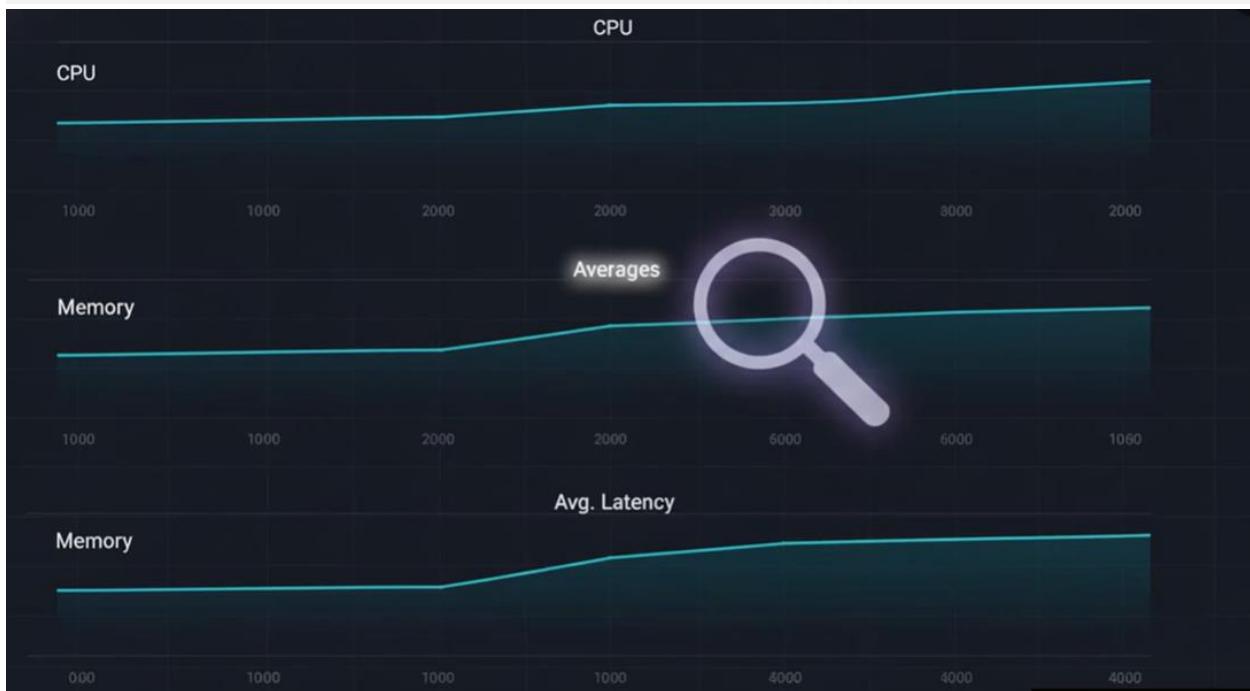


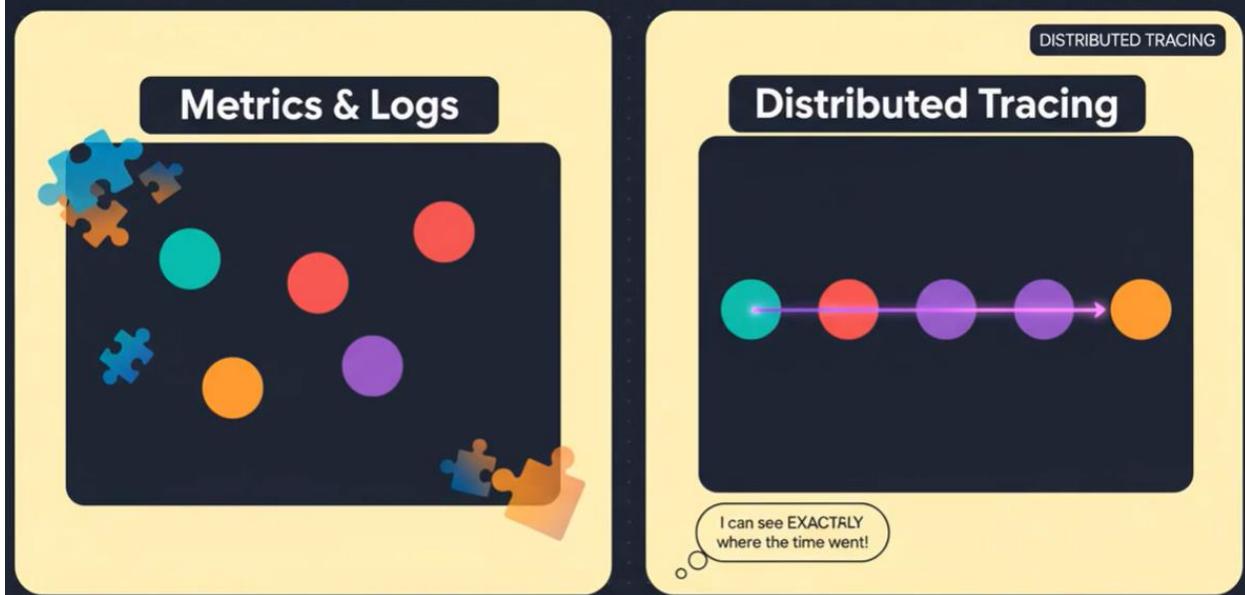
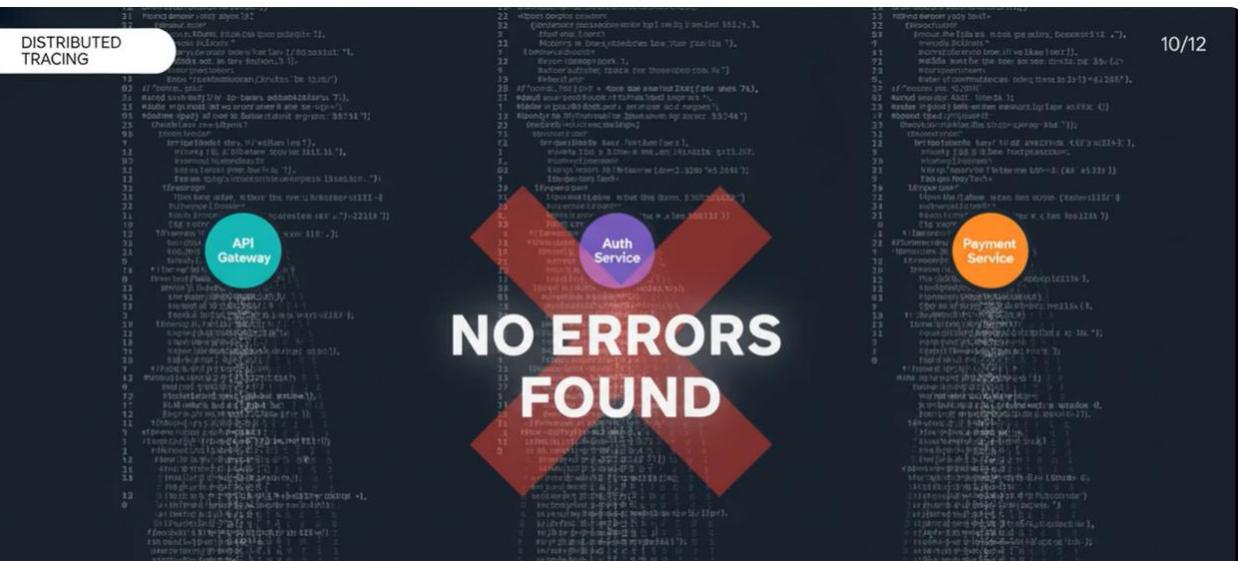
Configure Retention

Lifecycle Management

You have logs. Can you
now trace a request
across your system?

WHERE did those 3
seconds go?







A large yellow rectangular area contains the word 'Trace' in purple. Below it is a definition: 'The big picture of what happens when a request is made, representing its full path in your application.' To the right of the yellow area is a magnifying glass over a code snippet. The code shows a function call with parameters: 'span.start(): ((span.id()), span.start_end))'. At the bottom right is a legend for 'logs', 'events', and 'metrics'. The background features a grid and several small purple dots connected by arrows, representing a trace path. In the bottom right corner, the text 'Gheware DevOps AI' is visible.

Span

A single unit of work or operation within a trace, like an HTTP request. Spans are the building blocks.

DISTRIBUTED TRACING

10/12

350ms

API Call

DB Query

Cache Get

Process Data

HTTP POST /users (120ms)

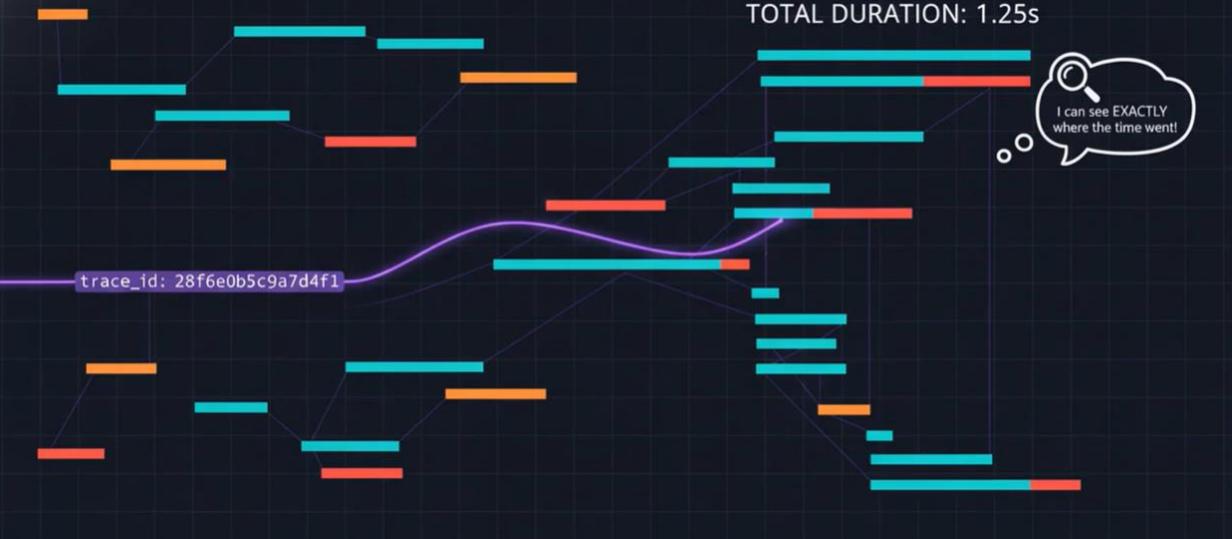


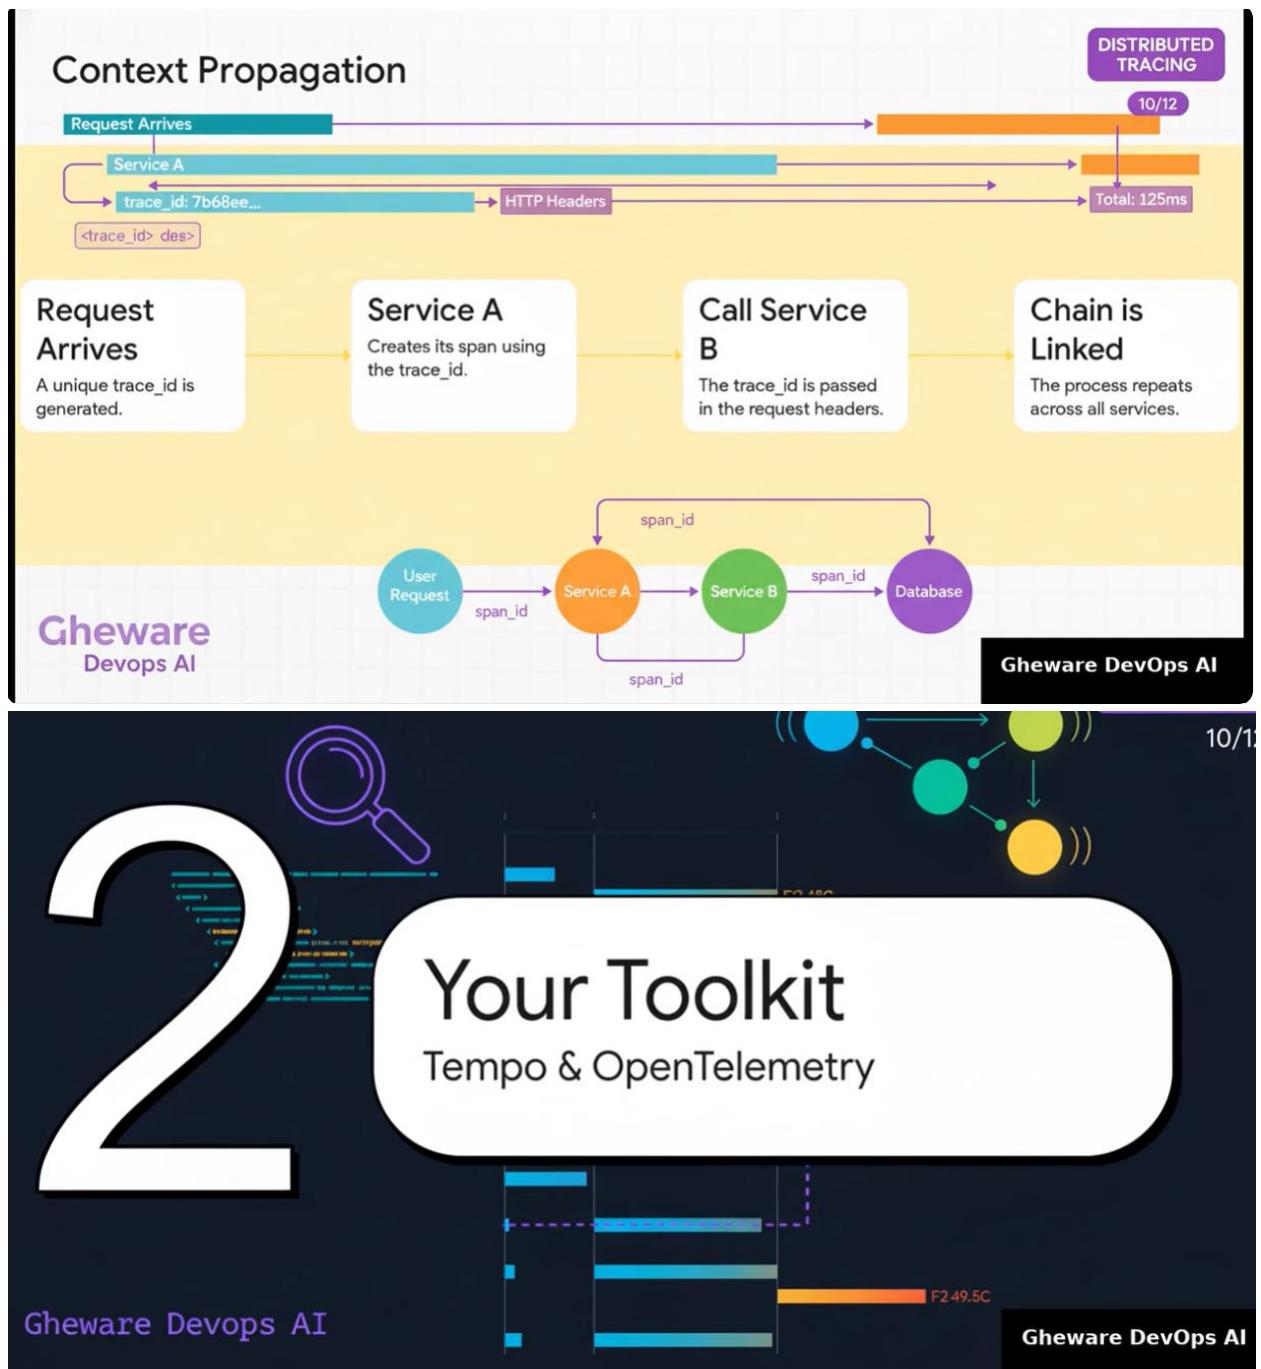
DISTRIBUTED TRACING

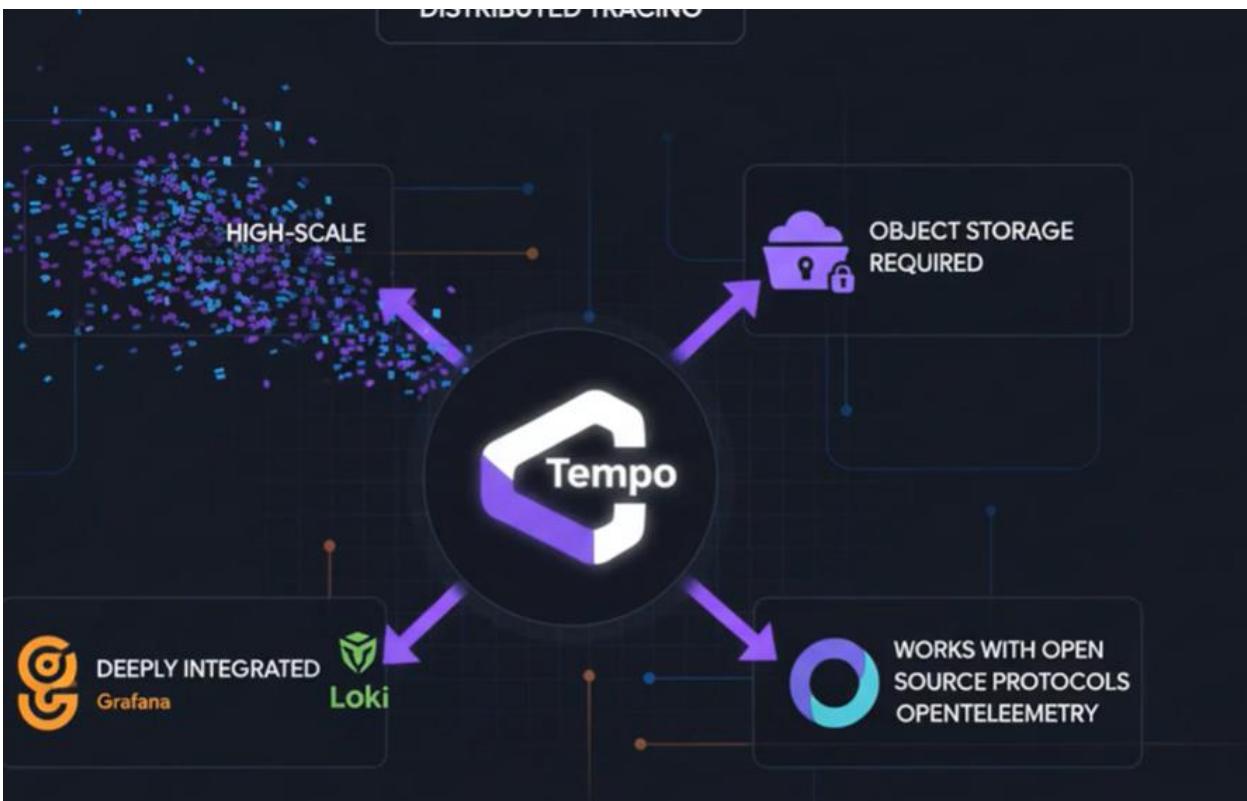
10/12

TOTAL DURATION: 1.25s

I can see EXACTLY where the time went!







Tempo is the backend, but your applications must first be instrumented to generate and emit traces.

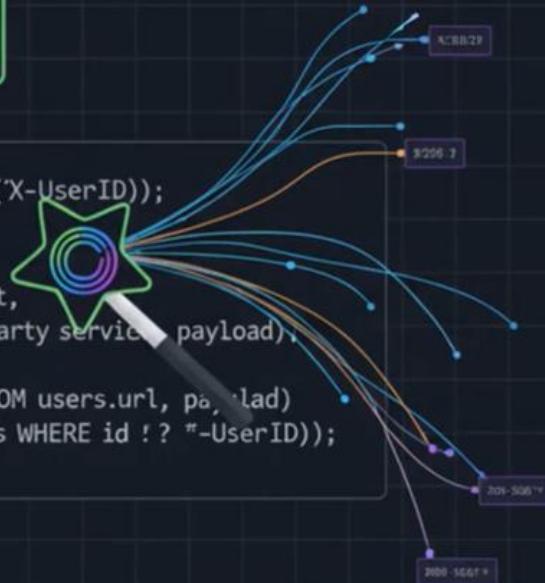
OpenTelemetry (OTel)

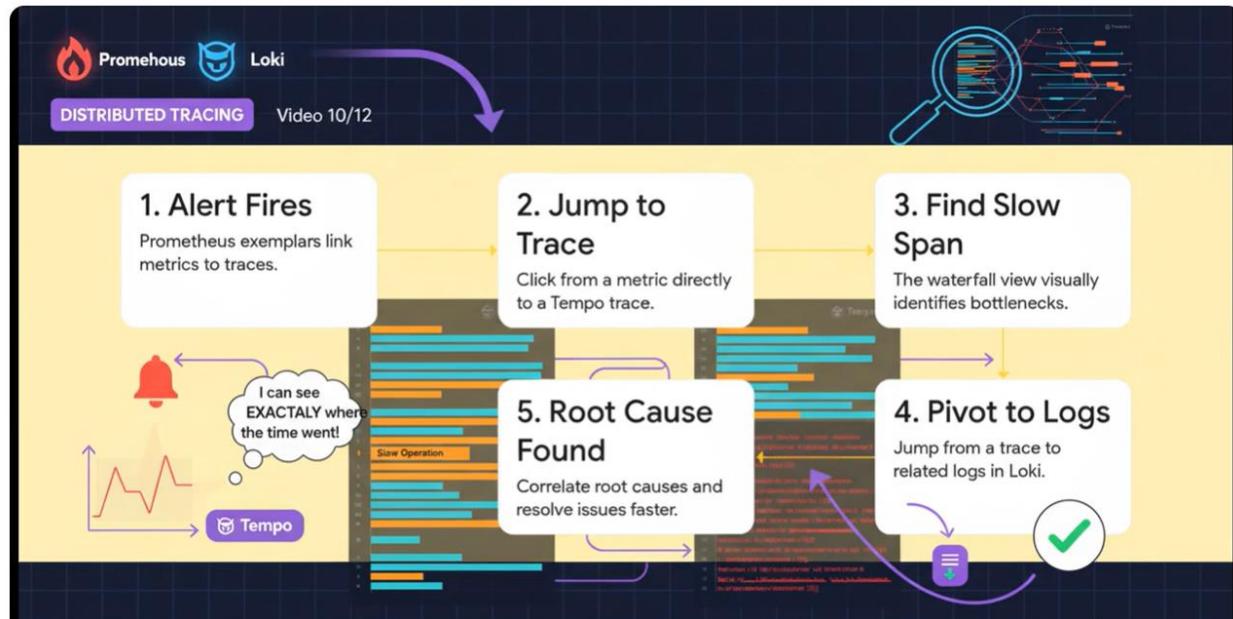
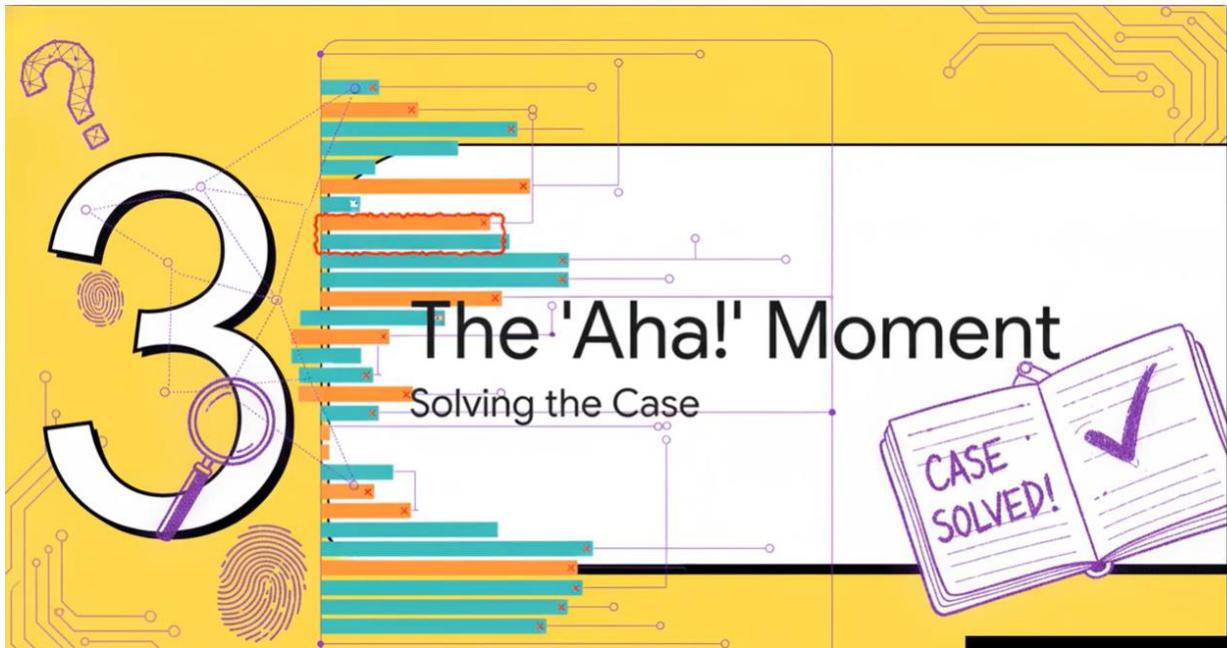
An observability framework to generate and collect application telemetry data like traces, metrics, and logs.



**NO CODE CHANGES
REQUIRED**

```
user-id = request.headers.get('X-UserID');
{
  response = http.client.post(
    http.client.post.third-party service, payload);
  response = (SELECT * FROM users.url, pa, glad)
    db.query((SELECT * users WHERE id != ? "-UserID"));
};
trace_id: 7b68ee...f5d4
```





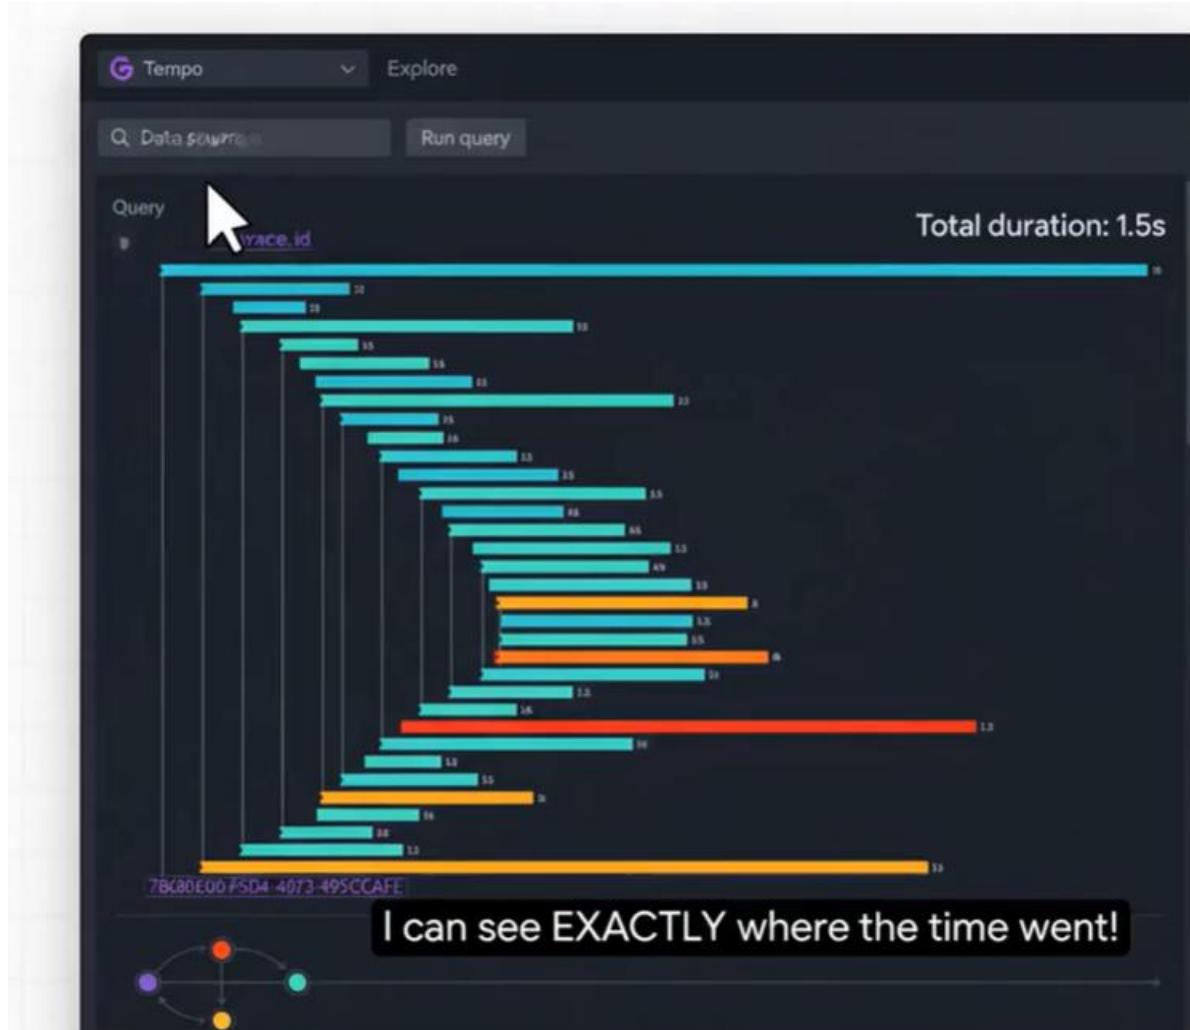
DISTRIBUTED TRACING

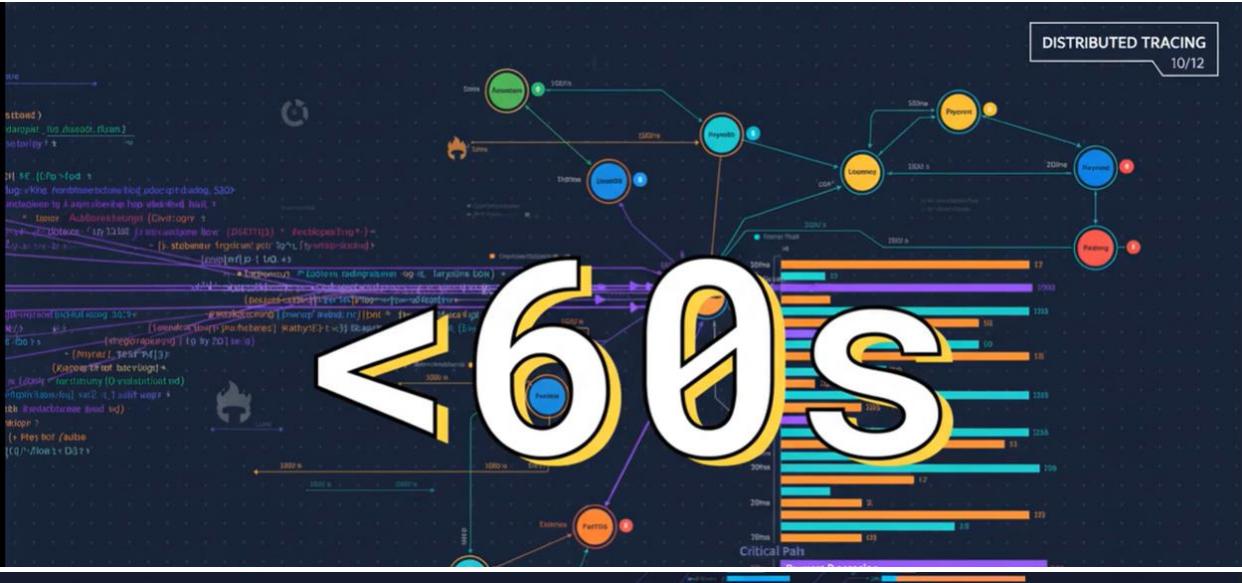
Tempo
Sprint
Destitutio
Refuge
Grottaatt
Souttamtors

```
file: calculatingIssue_a.wave.dts"
21  date_t, file_name_dts, share_traced_id, file_id, file_type, file_code );
22  curlload_file_filter = curl 47785648;
23  trace_id:484d((br:14, time_wanted,-4 ))
24;
25;
26  trace_id:484d((br:14, time_wanted,-4 ))
27  trace_id:484d((br:14, time_wanted,-4 ))
28;
29  trace_id:484d((br:14, time_wanted,-4 ))
30  trace_id:484d((br:14, time_wanted,-4 ))
31;
32  trace_id:484d((br:14, time_wanted,-4 ))
33  trace_id:484d((br:14, time_wanted,-4 ))
34;
35;
36  trace_id:484d((br:14, time_wanted,-4 ))
37  trace_id:484d((br:14, time_wanted,-4 ))
38;
39  trace_id:484d((br:14, time_wanted,-4 ))
40  trace_id:484d((br:14, time_wanted,-4 ))
41;
42  trace_id:484d((br:14, time_wanted,-4 ))
43  trace_id:484d((br:14, time_wanted,-4 ))
44;
45  trace_id:484d((br:14, time_wanted,-4 ))
46  trace_id:484d((br:14, time_wanted,-4 ))
47;
48  trace_id:484d((br:14, time_wanted,-4 ))
49  trace_id:484d((br:14, time_wanted,-4 ))
50;
51  trace_id:484d((br:14, time_wanted,-4 ))
52  trace_id:484d((br:14, time_wanted,-4 ))
53;
54  trace_id:484d((br:14, time_wanted,-4 ))
55  trace_id:484d((br:14, time_wanted,-4 ))
56;
57  trace_id:484d((br:14, time_wanted,-4 ))
58  trace_id:484d((br:14, time_wanted,-4 ))
59;
59-
```

trace_id=alb2c34e5f6 —  key to unlocking the full request journey

```
32  trace_id:484d((br:14, time_wanted,-4 ))
33;
34;
35;
36;
37;
38;
39  trace_id:484d((br:14, time_wanted,-4 ))
40  trace_id:484d((br:14, time_wanted,-4 ))
41;
42;
43;
44;
45;
46;
47;
48;
49;
49-
```





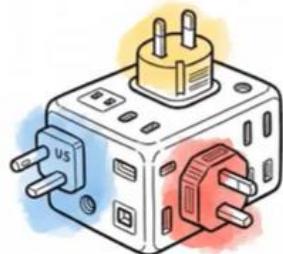


Your Takeaway

- Tracing shows the full request journey.
 - Spans pinpoint individual operations.
 - OpenTelemetry is the instrumentation standard.
 - Grafana Tempo is the high-scale backend.
 - Grafana provides integrated visualization.

What hidden performance mysteries will you uncover?

OpenTelemetry Explainer



Your application is **slow**. Users are complaining. But which microservice is the problem?

1

Instrumentation Chaos

Before OpenTelemetry

THE OLD WAY



Multiple, chaotic agents for each vendor.

THE NEW WAY



One standard, clean pipeline to any backend.

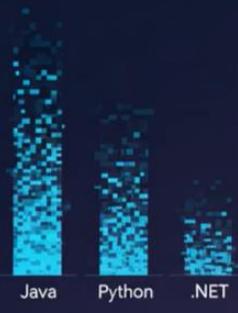
The Pain of Chaos



VENDOR LOCK-IN



HIGH COST



INCOSISTENT QUALITY



CNCF

OpenTelemetry

An open source, vendor-agnostic observability framework to generate, collect, and export telemetry data.

The Three Signals

- Traces: The path of a request through your application.
- Metrics: A measurement captured at runtime.
- Logs: A recording of an event.



A graduated project of the Cloud
Native Computing Foundation (CNCF)

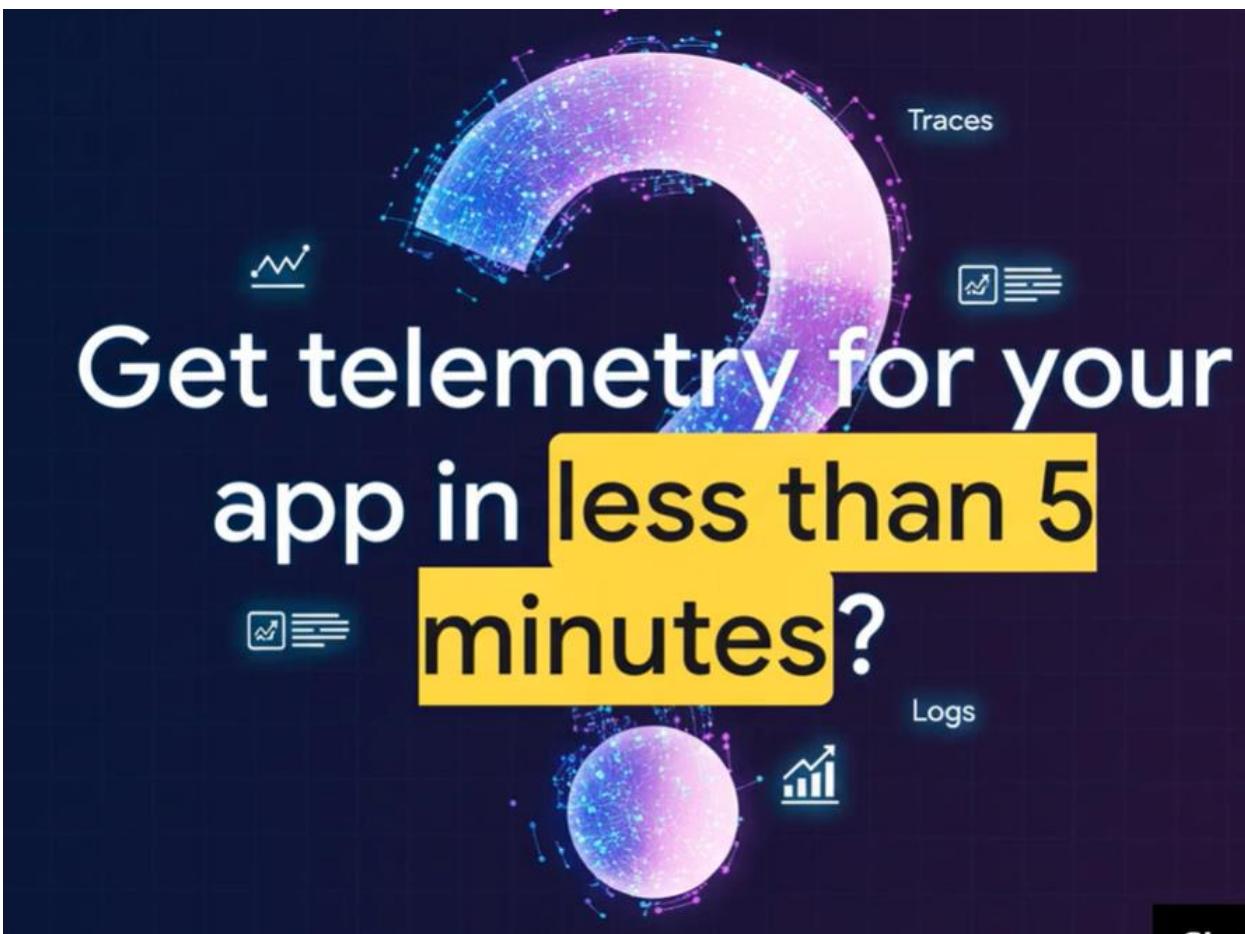


merging OpenTracing and OpenCensus to create
one powerful standard.

3

Instrumentation in Action

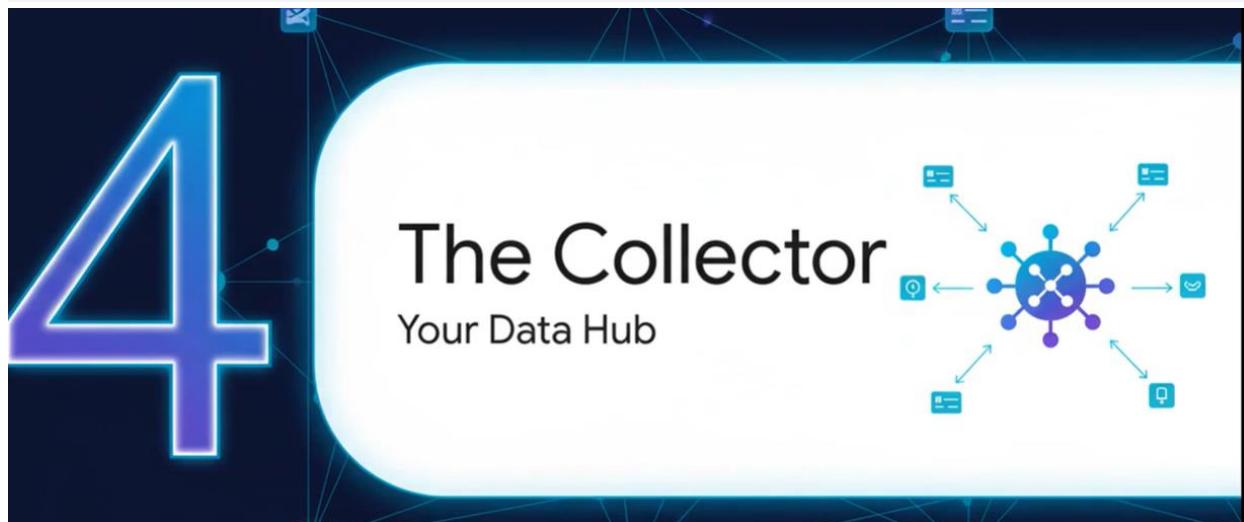
How It Works



Auto-Instrumentation



For more detail, use the API to add **custom** business context like a customer ID or order value, creating **richer** telemetry.



The OTEL Collector

A vendor-agnostic proxy that offers a unified way to receive, process, and export telemetry data.

Collector Pipeline

1. Receivers

Get data in various formats like OTLP, Jaeger, or Prometheus.

2. Processors

Modify data through batching, filtering, or sampling.

3. Exporters

Send processed data to one or more backend systems.

Why Use a Collector?

- Cost Control: Sample data to reduce backend costs.
- Reliability: Buffer data and retry sends to prevent data loss.
- Flexibility: Route data to multiple backends simultaneously.

One standard for all services.
One format for all your data
(traces, metrics, logs). One
flexible pipeline to manage it
all.



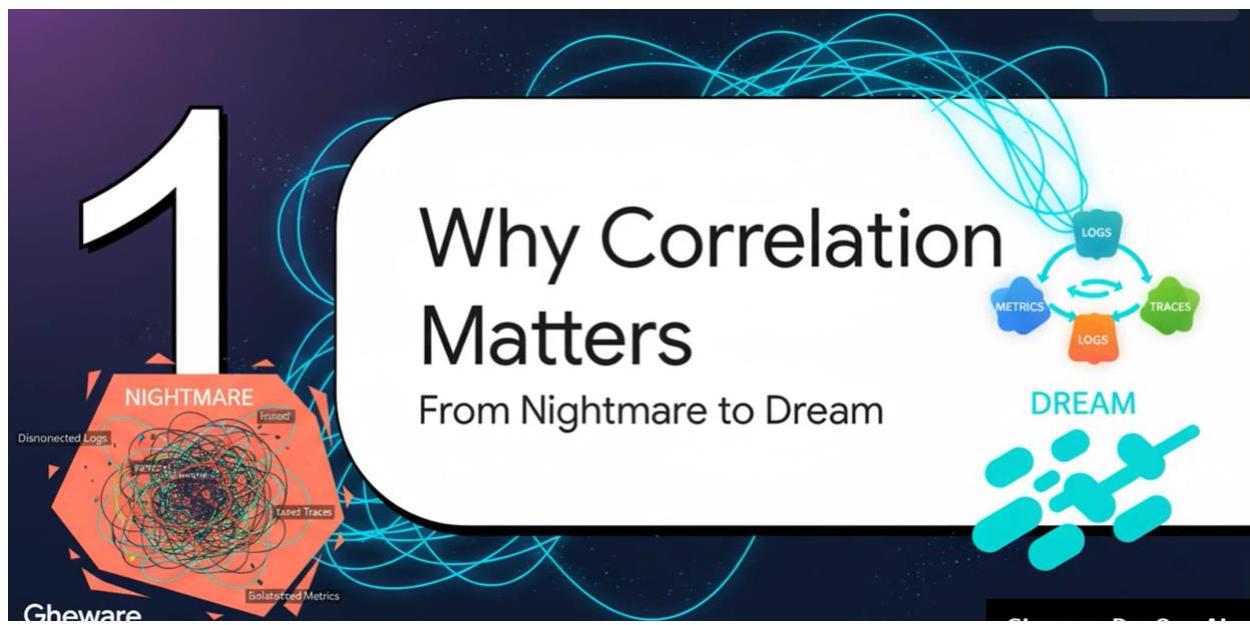
**Your application has a
story to tell. Are you
listening?**

Correlating Metrics, Logs, Traces



Ghewa

An alert fires. **WHERE** is
the problem? **WHICH**
request failed?



GHEWARE DEVOPS AI
THE ULTIMATE OBSERVABILITY PLAYBOOK – SERIES FINALE

THE NIGHTMARE

Manual Debugging: 30+ Minutes of Frustration



30+ minutes

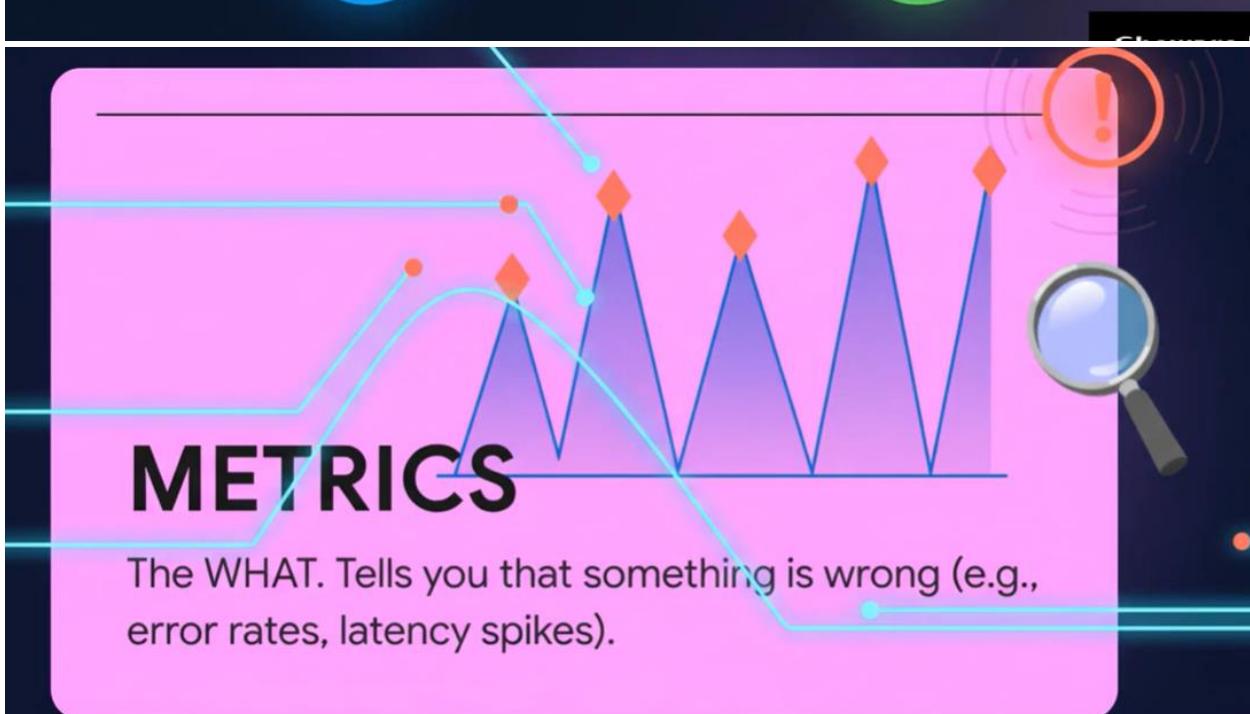
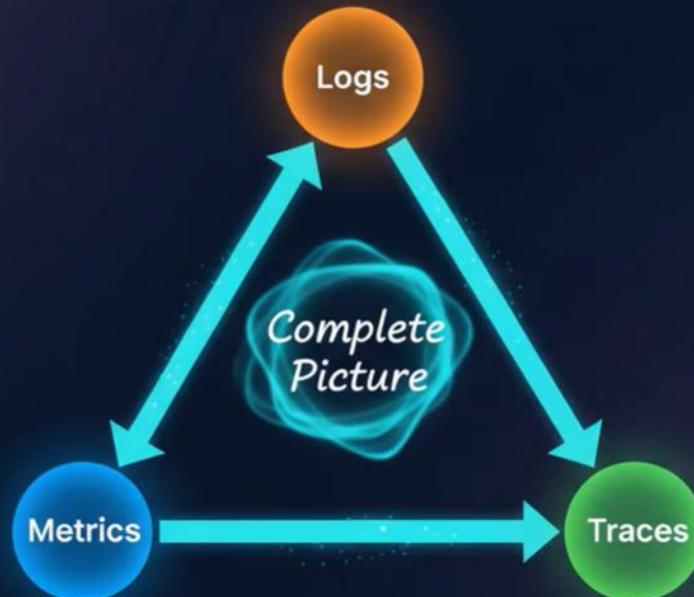
THE DREAM

Correlated Observability: Root Cause in < 2 Minutes.



< 2 minutes

Metrics, logs, and traces are no longer siloed. They are linked together, telling a single cohesive story.



```

51 trace:create msg="connection timbdout"
52   service=97}
53
54 timestamp=... 6==3) {
55   timestamp=... level:EROR msg="Connection timbdout",
56   service=4)
57   trace:definition failed for retien: snobits:database
58   error:code=104 commet=xyz
59
60 timestamp=... level:EROR for ser=database
61   DNS resolution failed for sobjem (e., "Connection titbase")
62   service=api
63
64 trace:iddia 503 == (=-t52
65   LENS resolution failed for servion ong= service=api
66   error:code=104 connection service=xyz
67   auth-service:connection sevlecid, monecculaturan))
68   in: 1000 ms. The connection was closed.
69   reason:timedout
70
71 timestamp=... {
72   tryave=... level:EROR for useren cretion database
73   service=4)
74   error:code=101
75   trown=10-08013=181
76   LENS resup= level:EROR conetion msg="Connection timcdout",
77   service=4))
78   (createtomap:

```

LOGS

The WHY. Provides context and error messages explaining the problem (e.g., "Connection timeout").

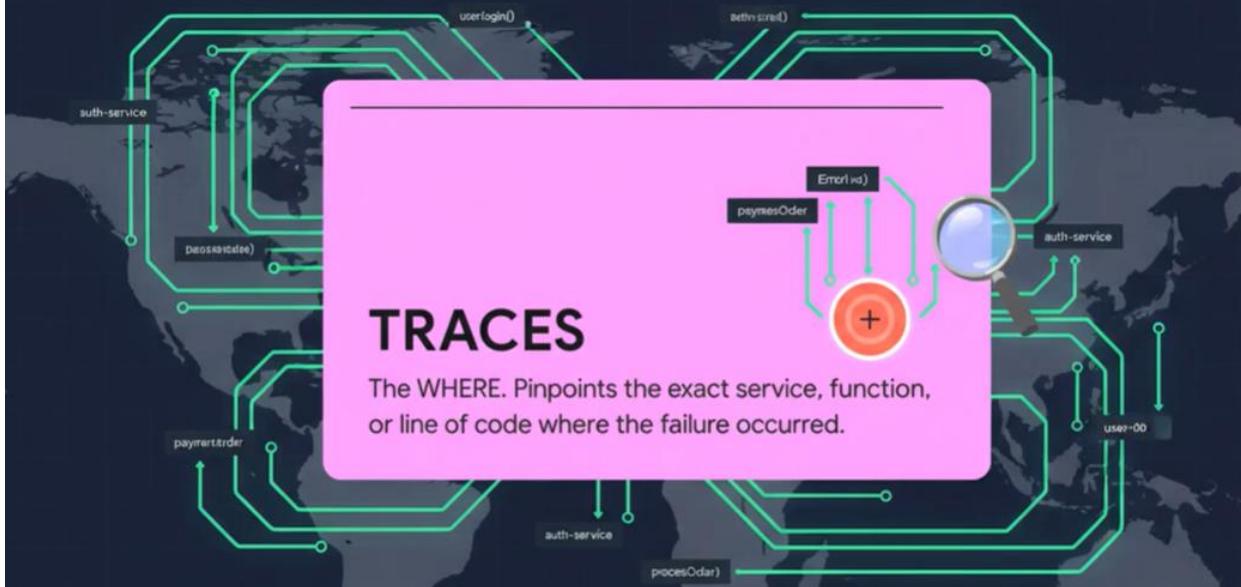
Series Finale

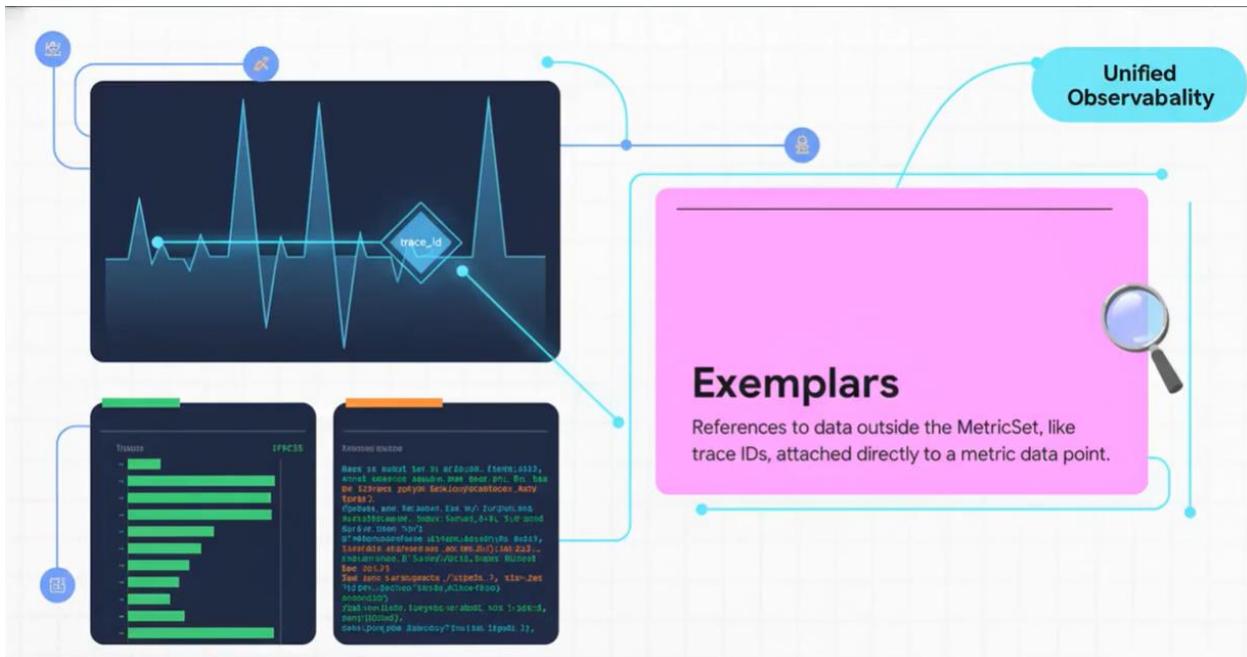
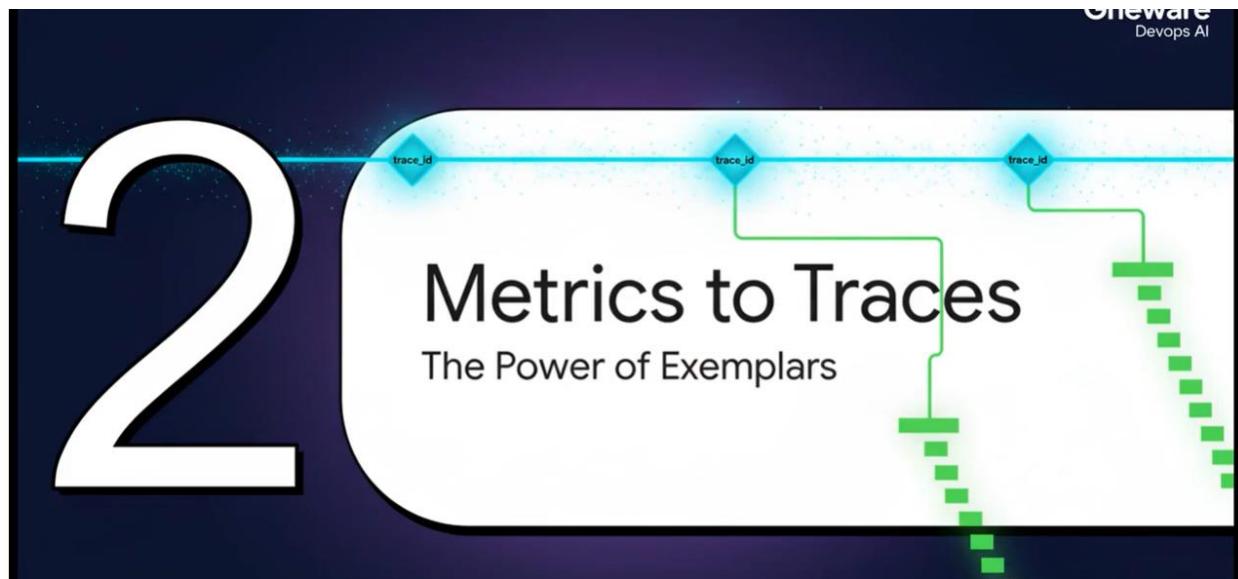
Gheware Devops AI

Gheware

TRACES

The WHERE. Pinpoints the exact service, function, or line of code where the failure occurred.





Enable this feature by adding the --enable-feature=exemplar-storage flag when starting your Prometheus server.



Logs to Traces
Clickable with Derived Fields

22CSE

Derived Fields

A feature to extract new fields from logs and create a link from the value using a regular expression.

Unified Observability

→ Tempo Trace View

Setting	Value
Name	TraceID
Regex	trace_id=(\w+)
Internal Link	On (select Tempo)
URL Label	trace_id

Logs

The image consists of two main sections. The top section is a white callout box containing text and a screenshot of a log viewer. The text reads: "Now, every trace_id in your logs is a blue, clickable link that opens the trace directly." Below the text is a screenshot of a log viewer titled "Split view". It shows a histogram on the left and a list of log entries on the right. One entry in the list has a blue underline over the "trace_id" field, indicating it is a clickable link. The bottom section is a large number "4" on the left, followed by a diagram in a rounded rectangle. The diagram illustrates a circular process between "Traces" and "Logs". A green arrow labeled "Trace" points from a trace icon to a log icon. A blue arrow labeled "Logs" points from a log icon back to a trace icon. The text "Completing the Circle" appears twice, once above the green arrow and once below the blue arrow.

“Now, every trace_id in your logs is a **blue, clickable link** that opens the trace directly.

4

Traces to Logs

Completing the Circle

Logs

Trace

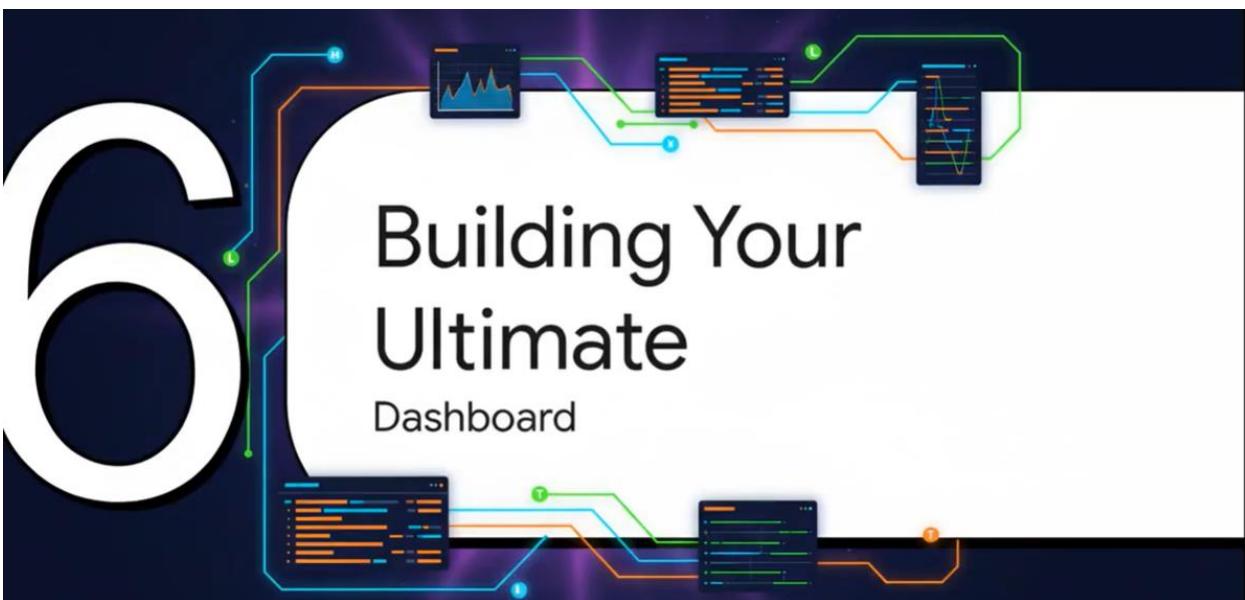
Completing the Circle

Tempo to Loki Config

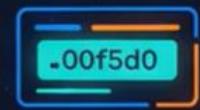


From Alert to Root Cause





Use consistent labels across all signals (metrics, logs, traces).



Ensure trace_id is included in ALL application logs.

Dashboard Best Practices



Enable exemplars on key golden signal histograms.



Test your correlations before an incident happens.

Journey Complete!

- Prometheus for metrics
- Loki for logs
- Tempo for traces
- OpenTelemetry for instrumentation
- Correlation for instant debugging!

